

# Keystone Annual Review 2023

Confidential Computing Consortium

Lily Sturmann [lsturman@redhat.com](mailto:lsturman@redhat.com)

Dayeol Lee [dayeolee@gmail.com](mailto:dayeolee@gmail.com)

# Goals of the Project

- ❑ Enable TEE on (almost) **all RISC-V processors**
  - Follow RISC-V standard ISA
  - Standard TEE specification for various RISC-V sub-ISA
- ❑ Make TEE **easy to customize** depending on needs
  - Base implementation vs. platform-specific implementation
  - Reuse the implementation across multiple platforms
- ❑ **Reduce the cost** of building TEE
  - Reduce hardware integration cost
  - Reduce verification cost
  - Integrate with existing software tools

# Remarks

## ❑ Code Maintenance

- Switched to [monorepo](#): for a better developer experience
- Bump [OpenSBI](#) v1.1

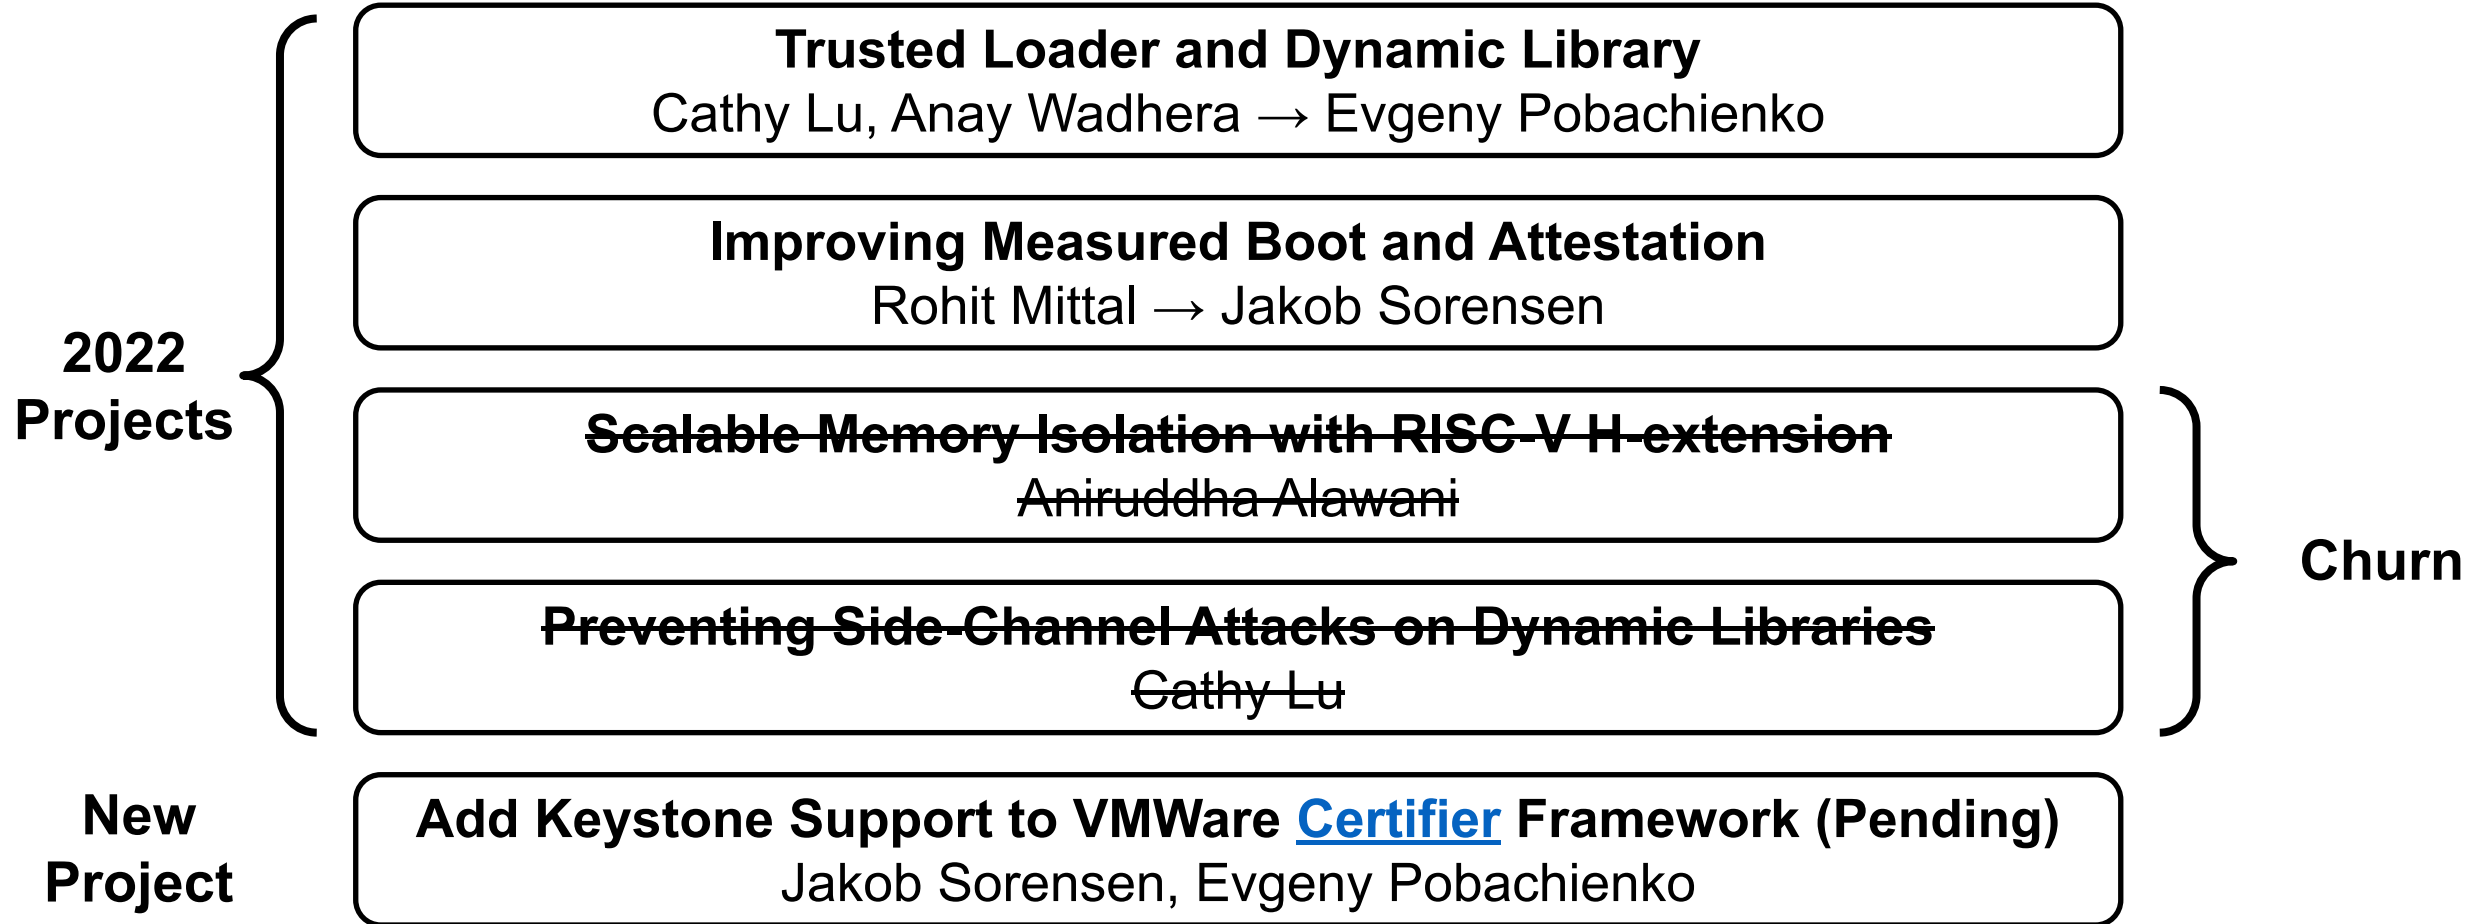
## ❑ The project have been very slow in 2022

- Five people from UCB graduated at the same time, and four of them left the project
- Less momentum from the industry

## ❑ Keystone is still a popular option in academia

- Gained 133 yearly citations (+28% YoY)
- 100+ forks mostly from researchers

# Subproject Status



# Why is the Project Stuck?

- ❑ Tight Coupling with RISC-V
  - Lack of Development Board
  - Many focused on low-end devices which is not Keystone is aiming for
  - RISC-V specification is still changing; no software standard yet
- ❑ Lack of Industry Contribution
  - Code quality geared toward research (not maintainability)
  - People leave the team after 1-2 years (usually at the same time)
- ❑ Lack of Application Demand
  - RISC-V software ecosystem is still growing, and the application demand is weak

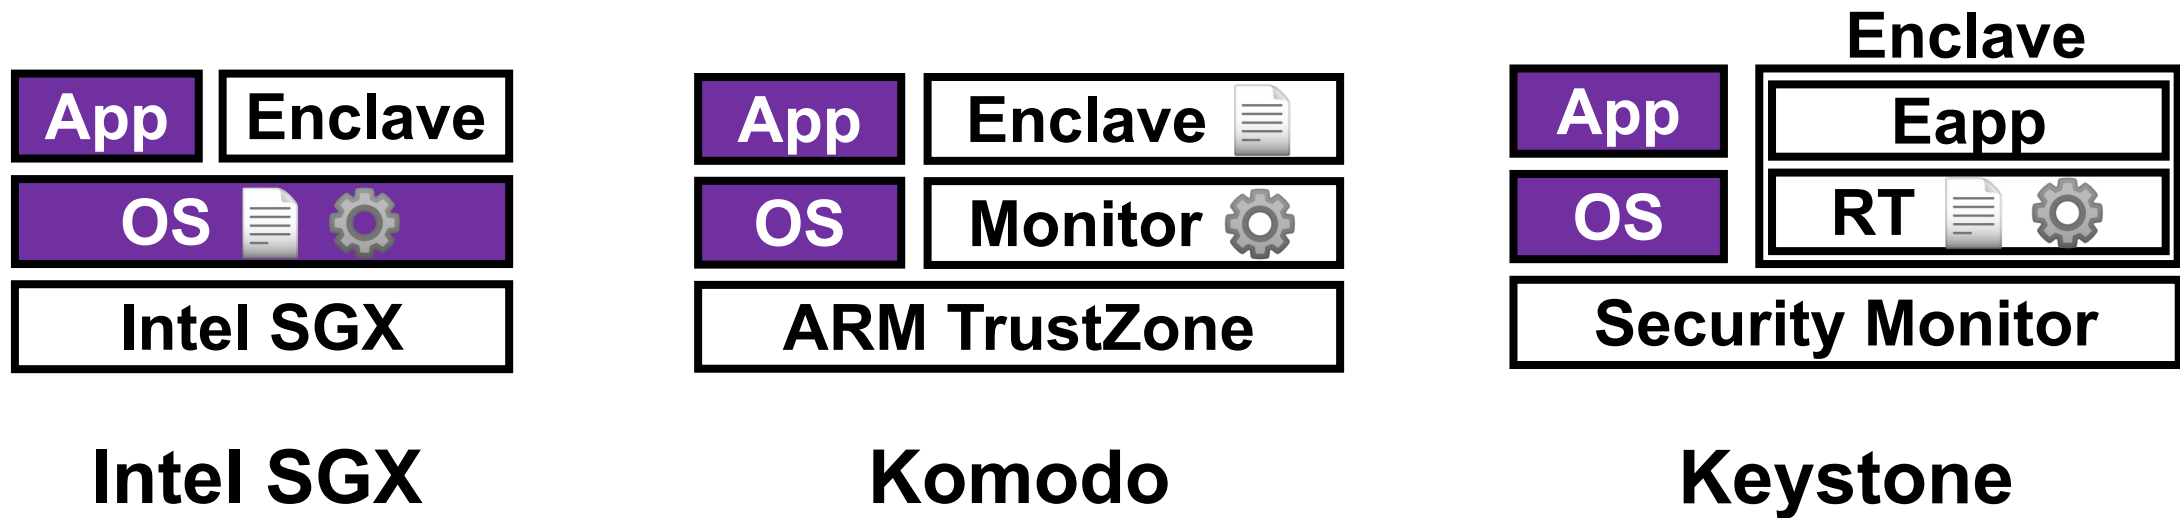
# Key Milestones for 2023

- Better application support
  - Dynamic library support
- Parity with industry standards
  - Standard crypto for measured boot / attestation
- Increase dev board accessibility
  - Participate in RISC-V development board program
  - Expecting a supply chain relief in mid 2023
- Work closely with RISC-V AP-TEE working group
  - Not directly relevant, but they are interested in pushing towards server-class RISC-V TEE in the future

# Thank You!

# Memory Management in Keystone

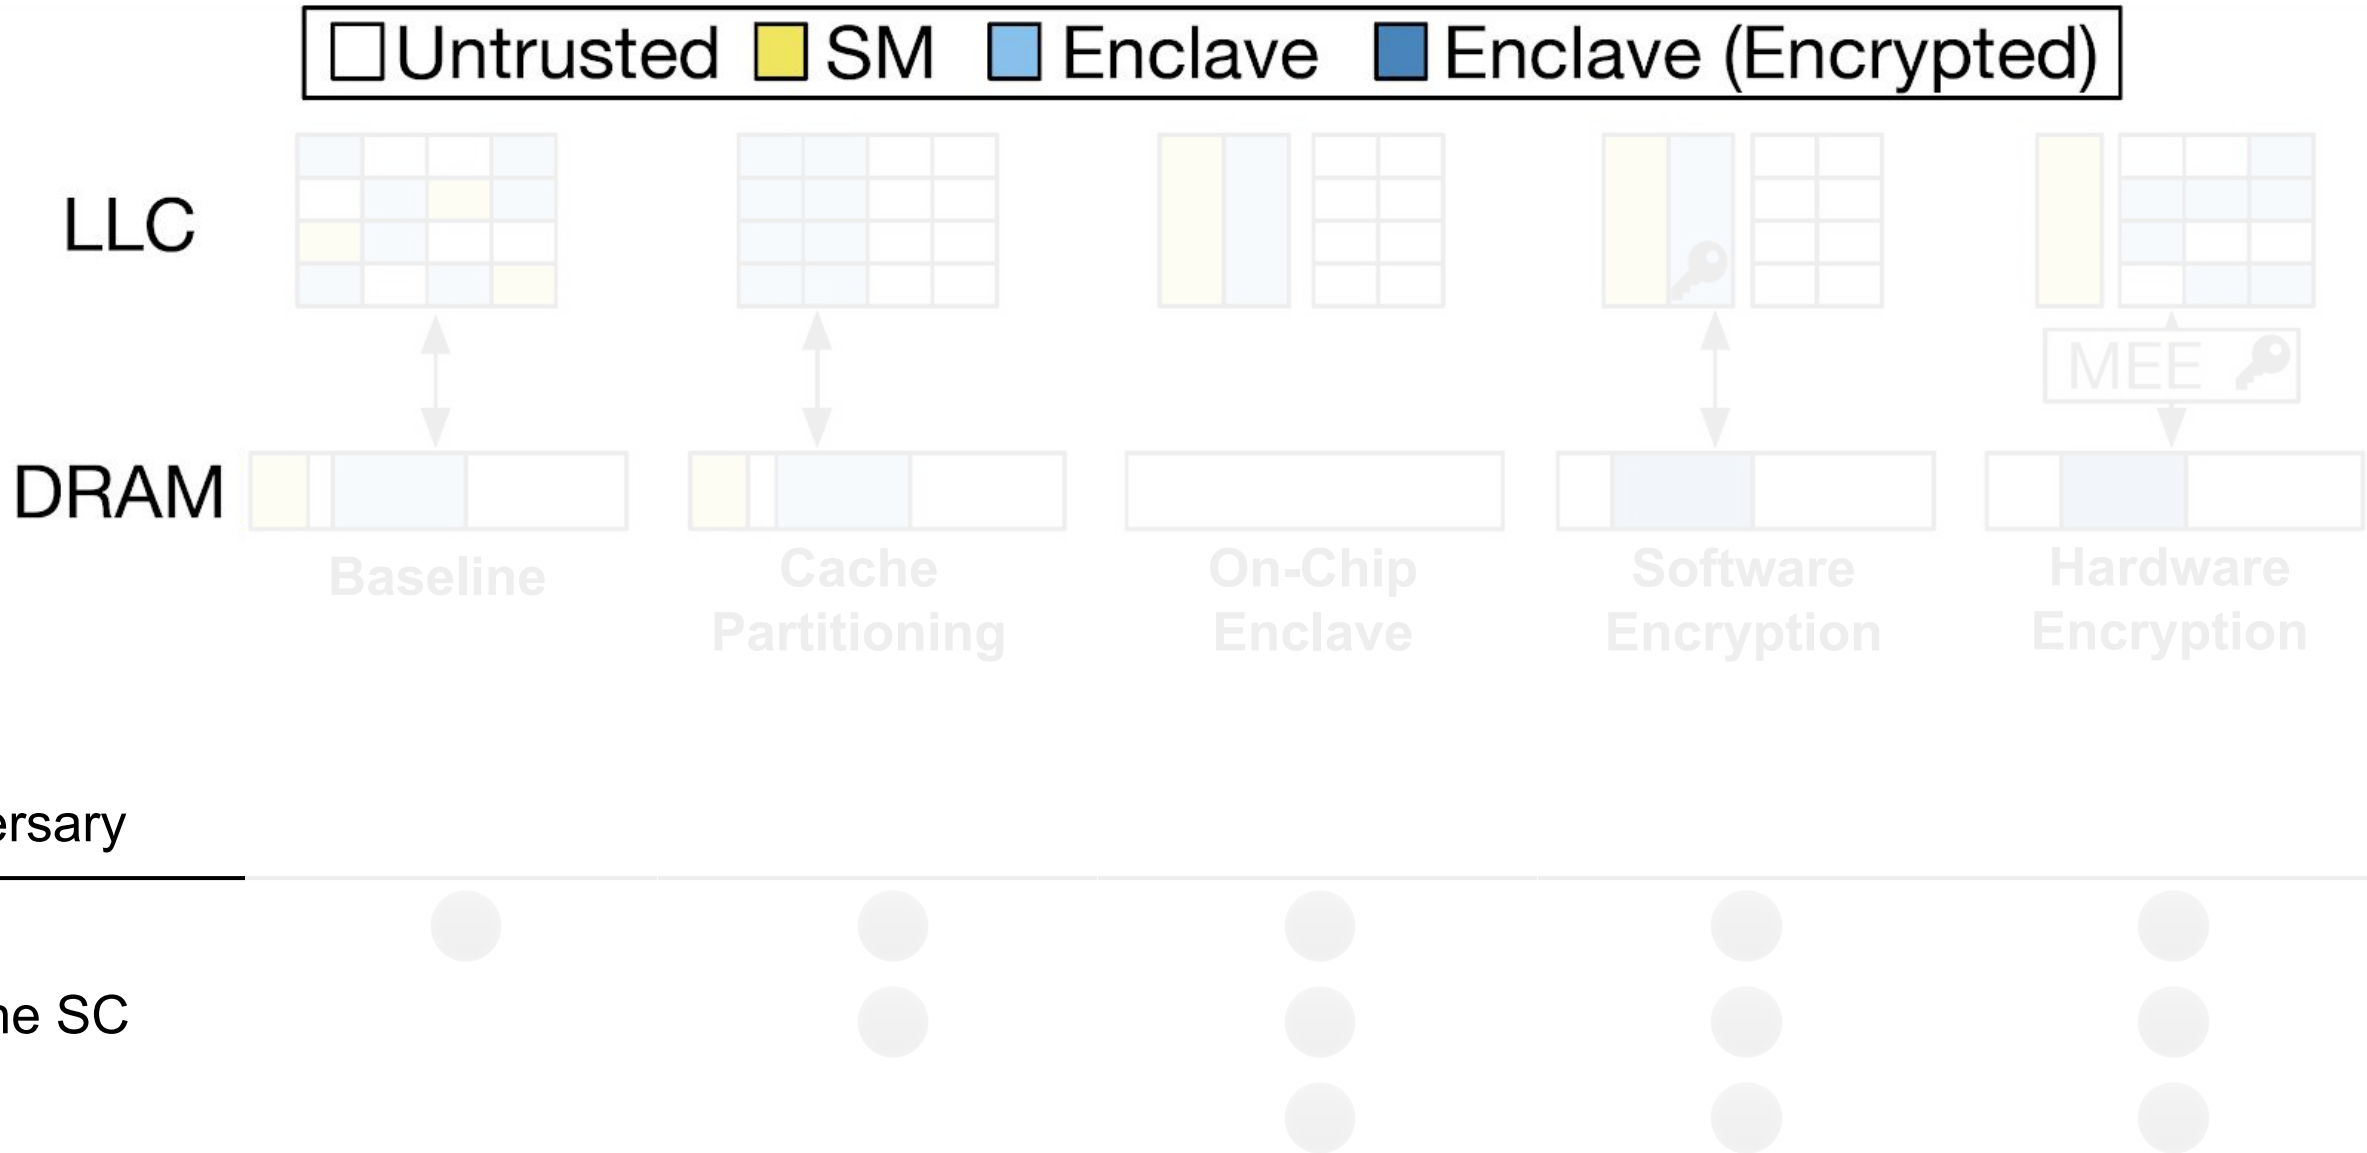
 = untrusted management  = page table  =



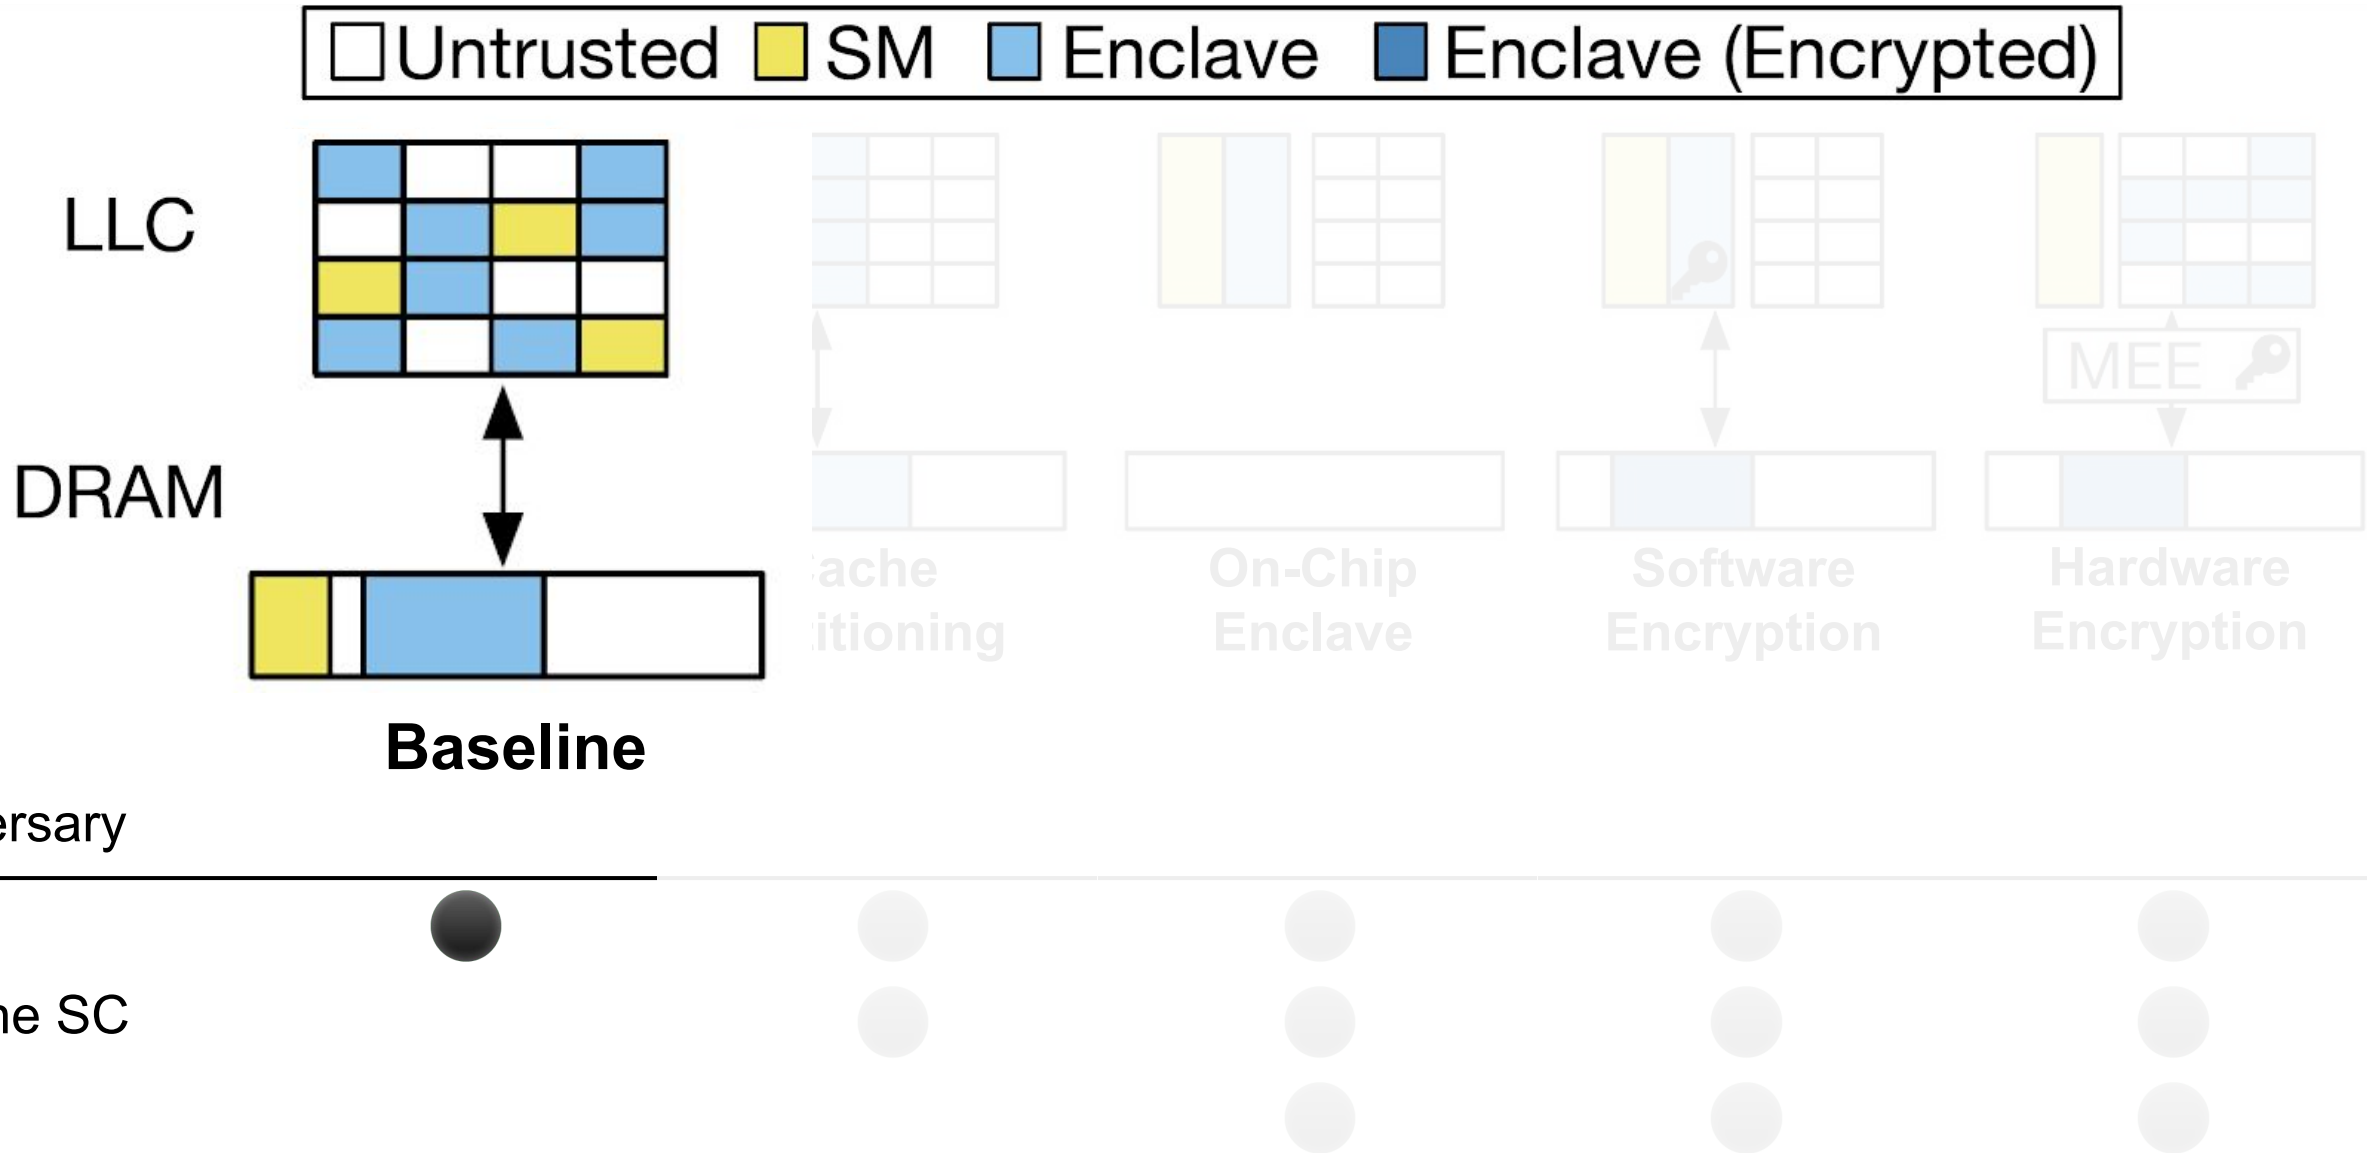
- ☐ Enclave self resource management (e.g., dynamic memory resizing)
- ☐ Various memory protection mechanisms



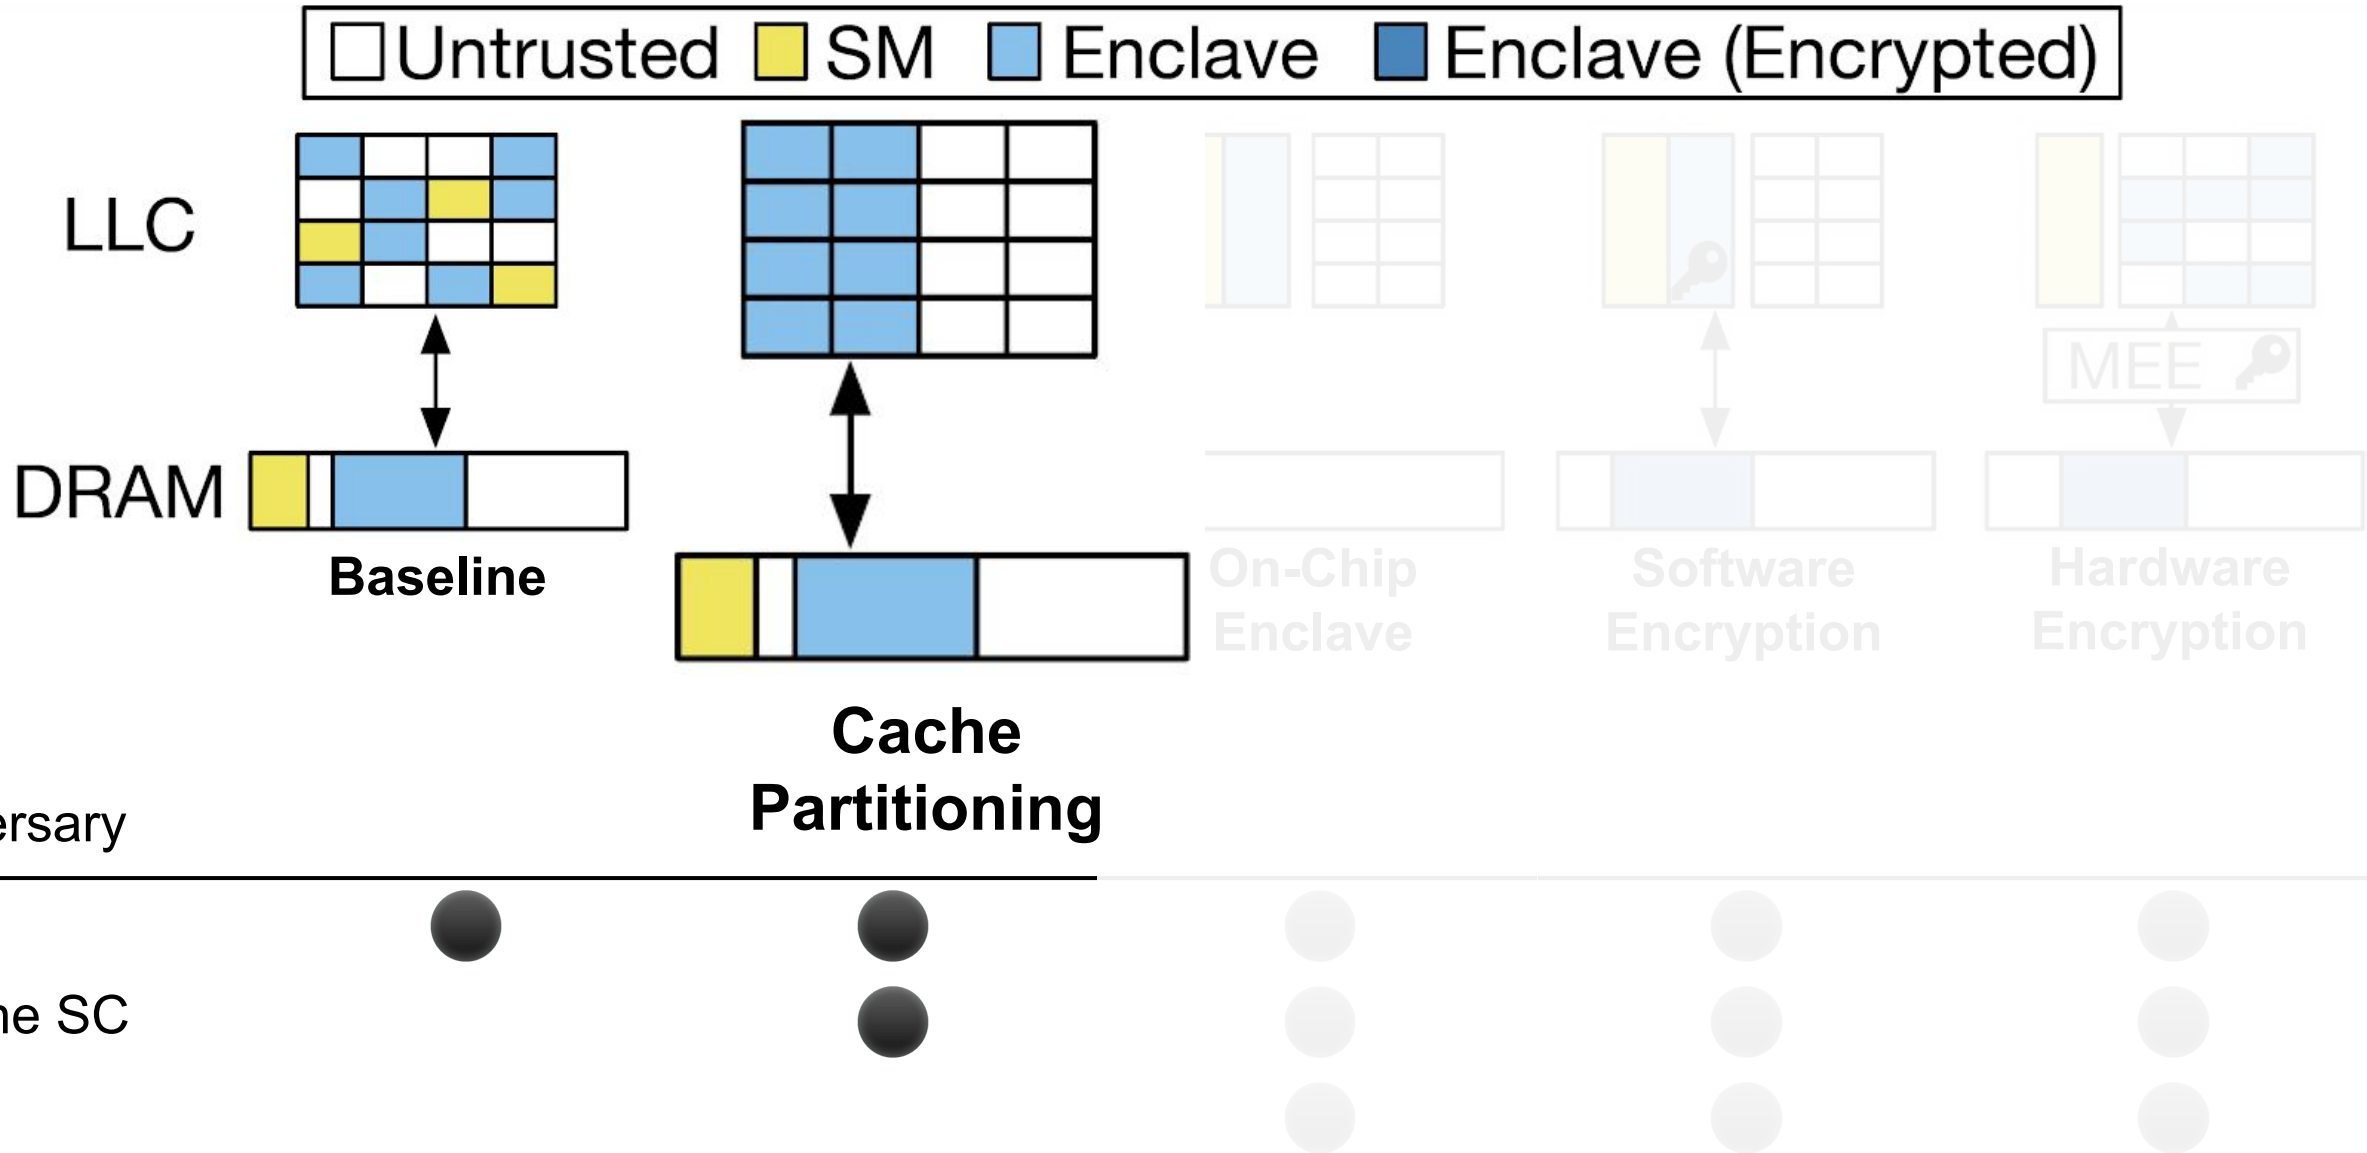
# Various Memory Protection Mechanisms



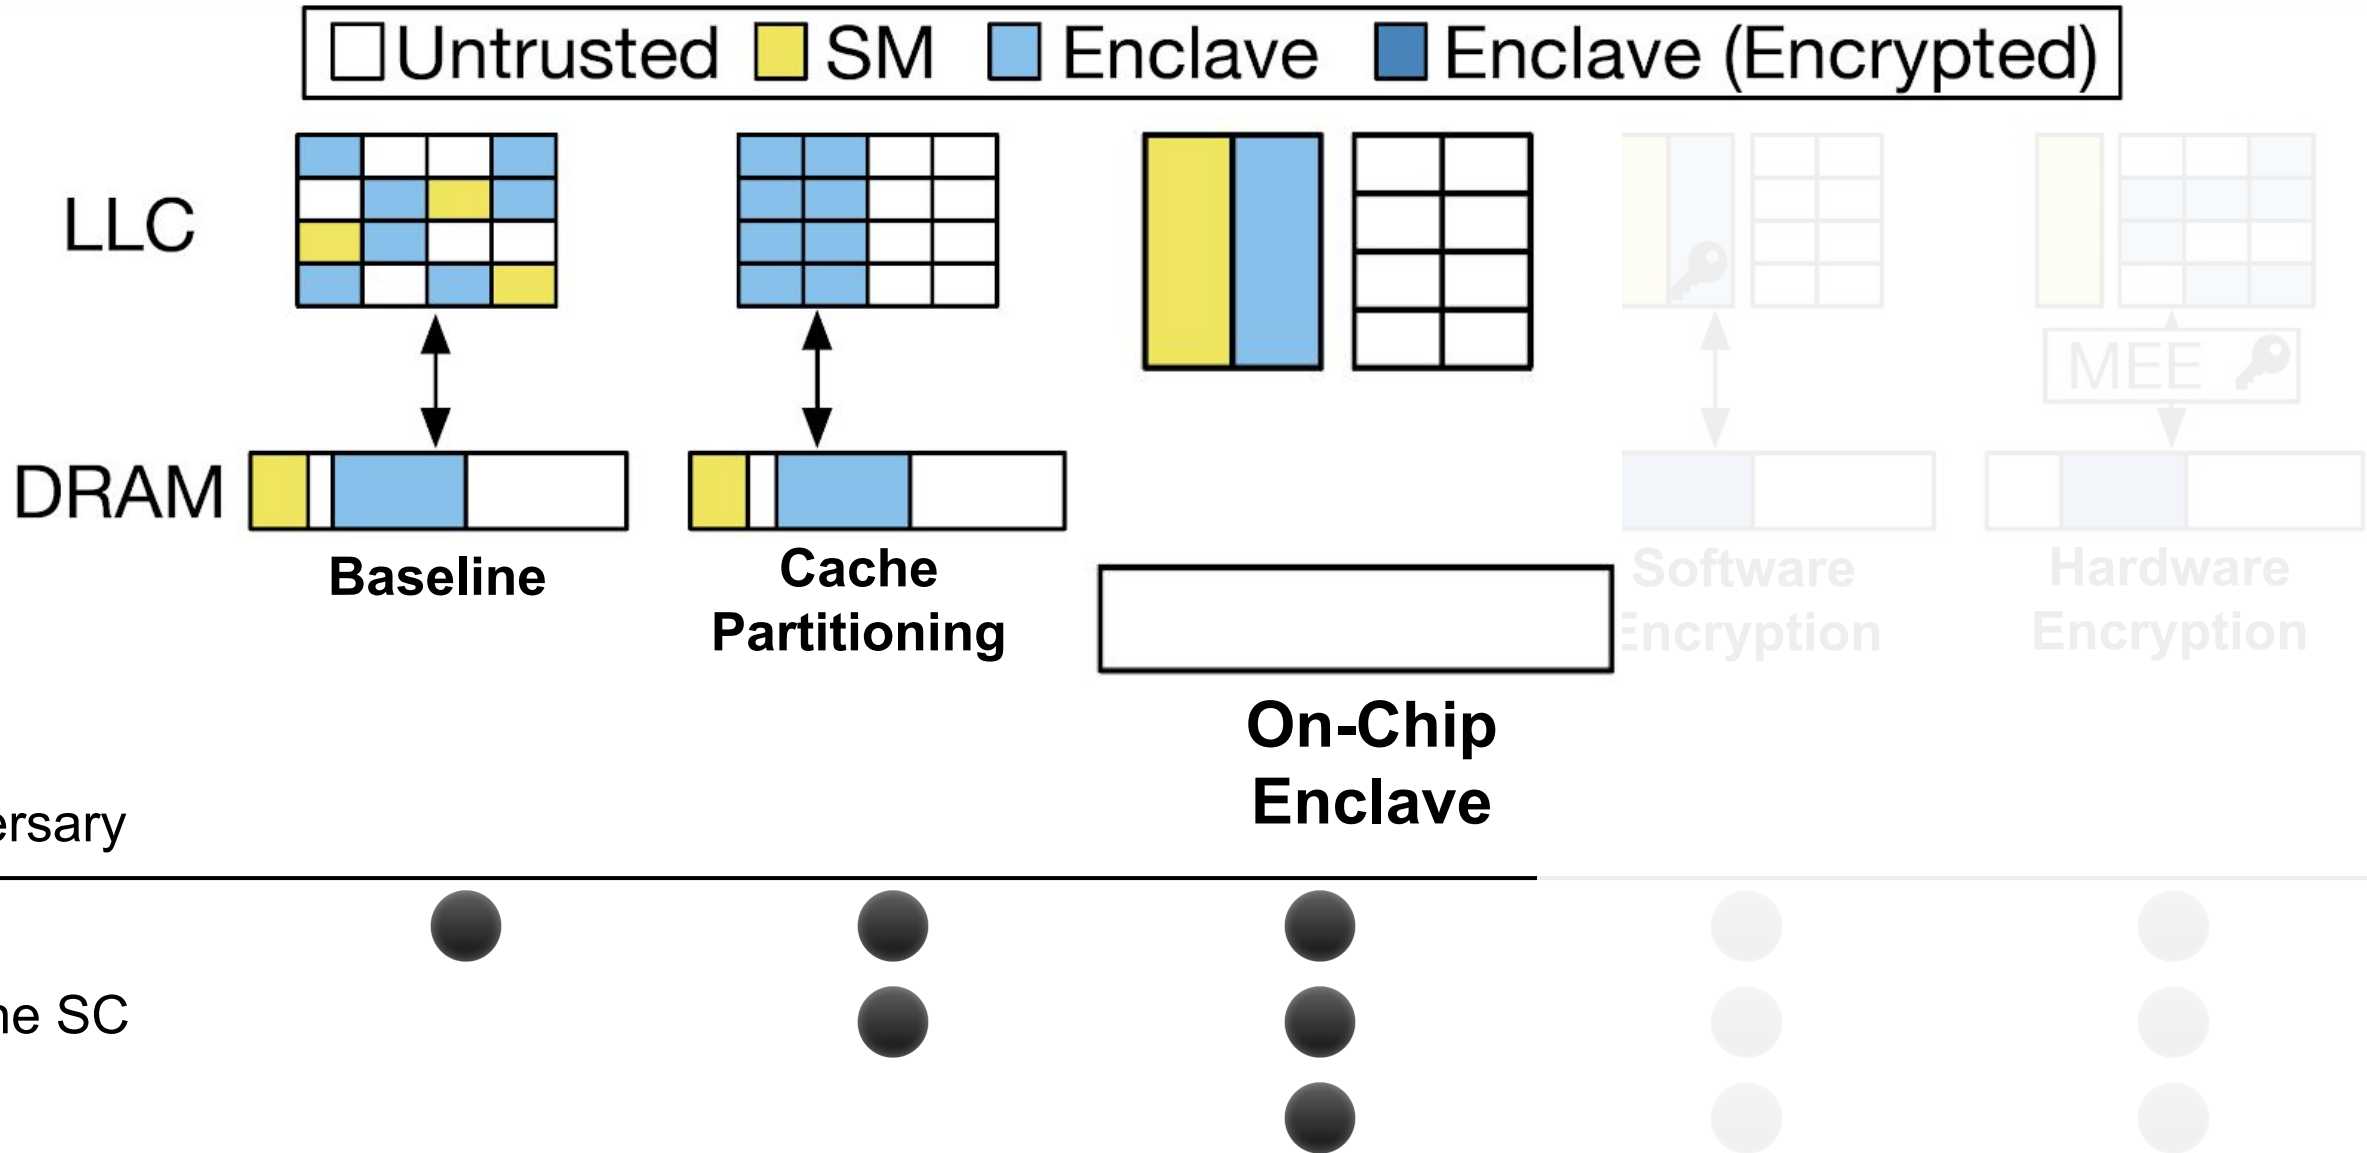
# Various Memory Protection Mechanisms



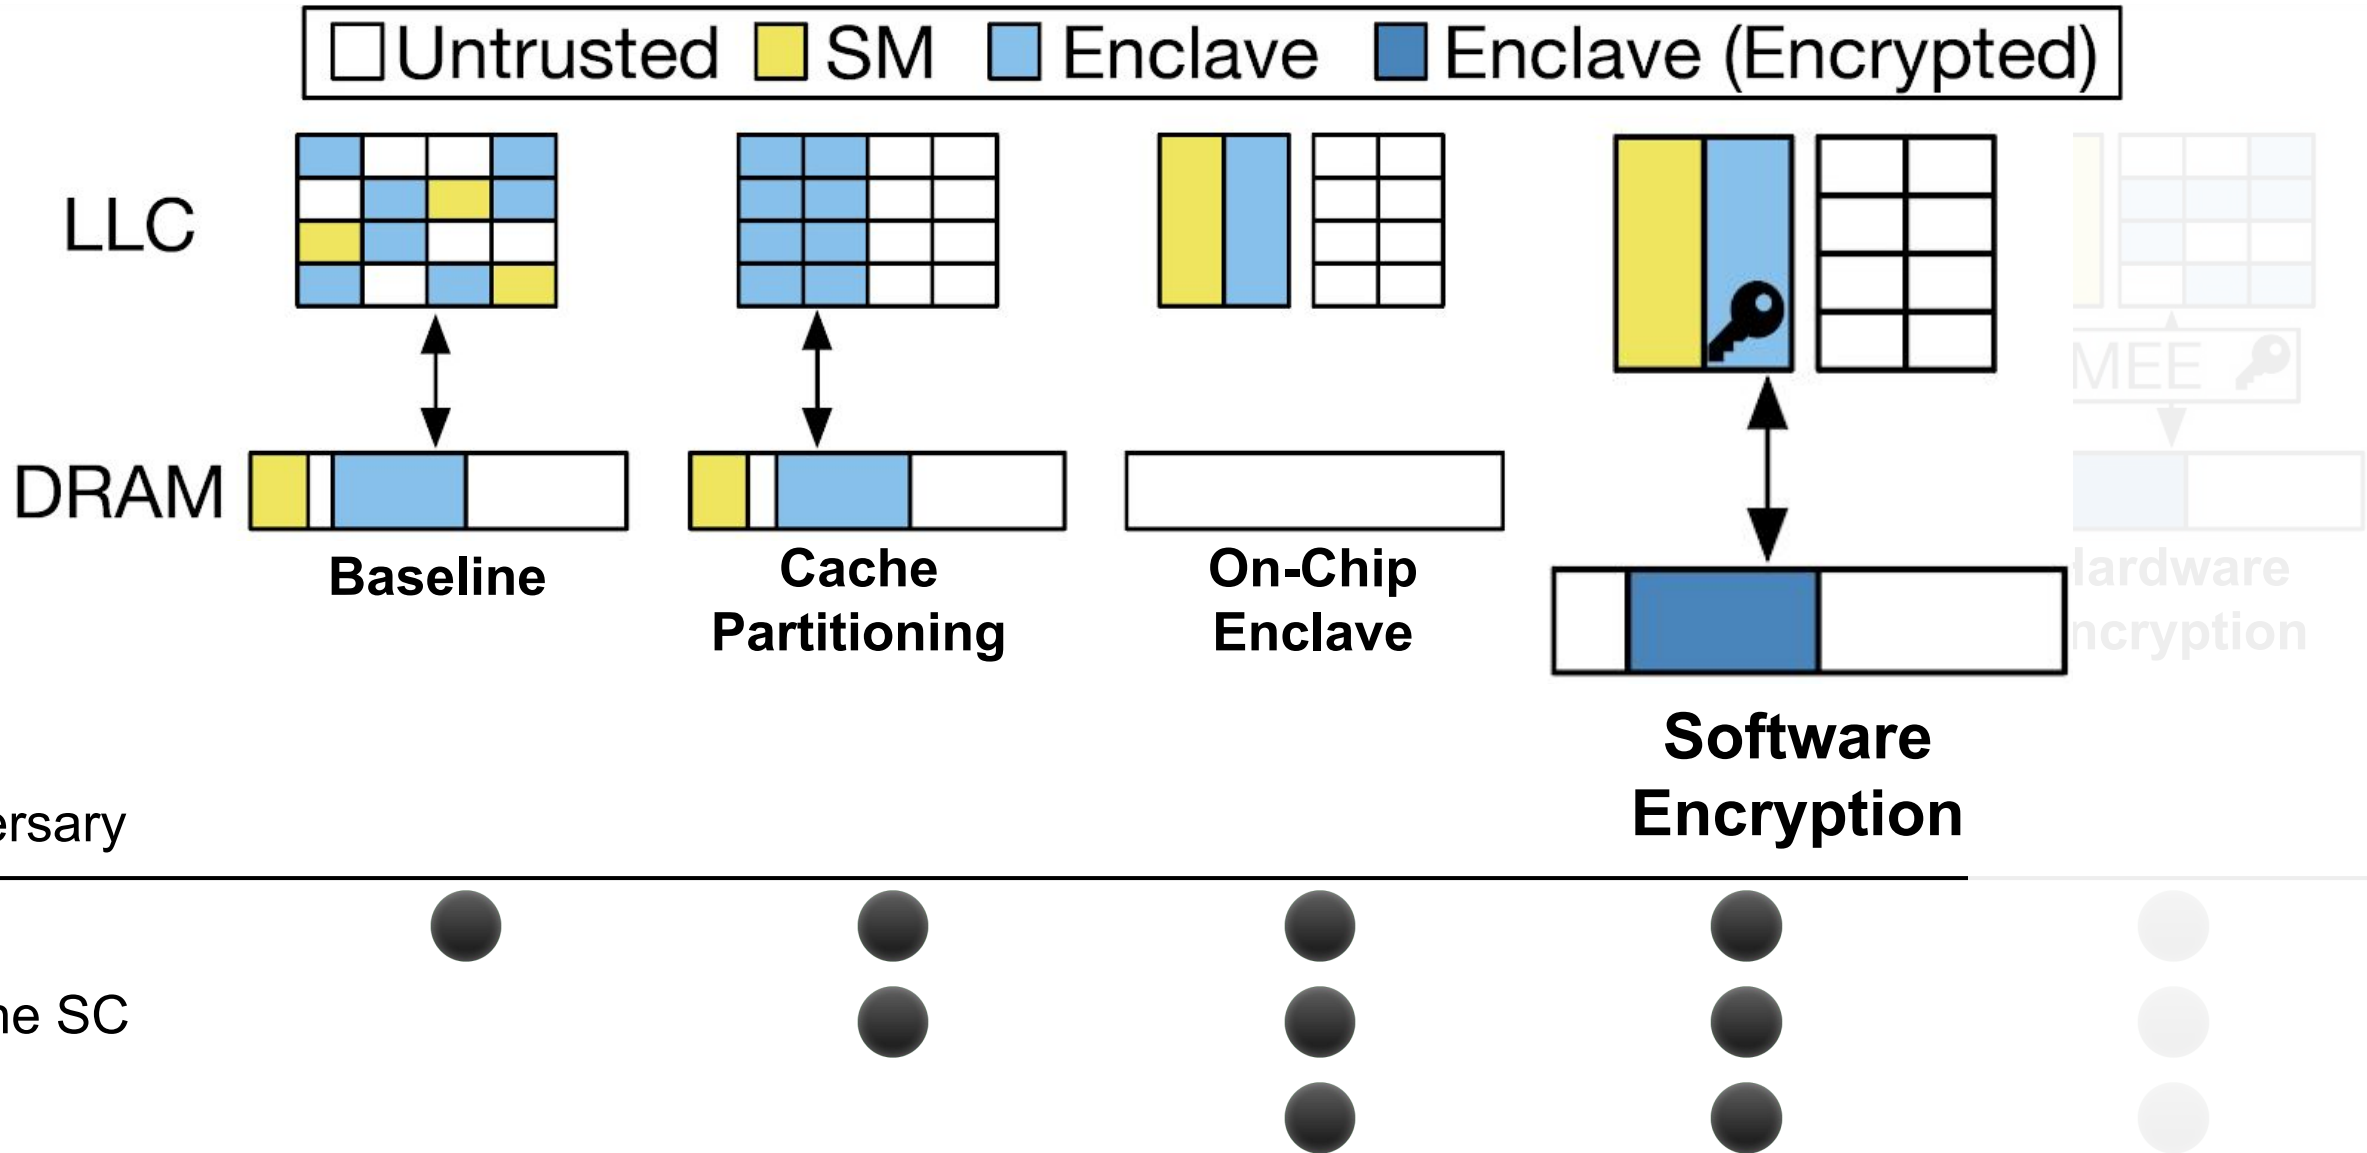
# Various Memory Protection Mechanisms




# Various Memory Protection Mechanisms



# Various Memory Protection Mechanisms



# Various Memory Protection Mechanisms



Adversary

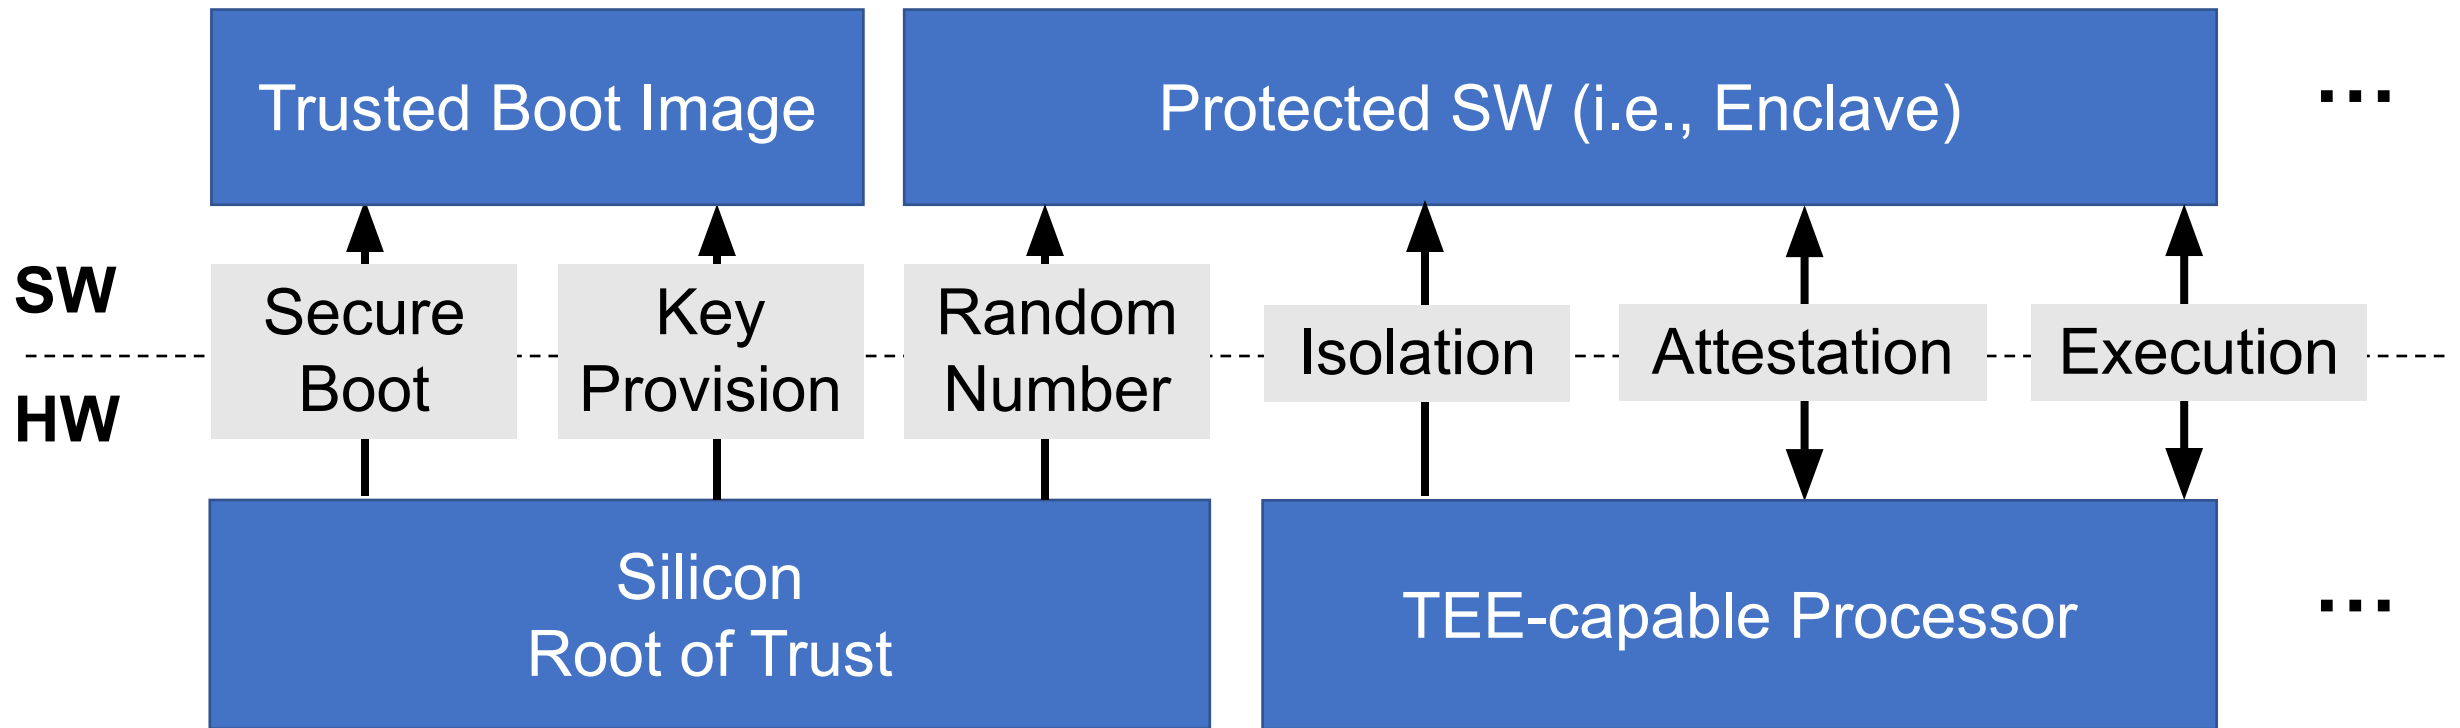
SW

Cache SC

HW




# TEE as Software-Hardware Contract



# Why Do We Need a Flexible Design?

- ❑ Trade-offs in security, functionality, and performance
- ❑ A reasonable threat model depends on
  - ❑ Different platforms (e.g., mobile vs. desktop)
  - ❑ Different applications (e.g., gemm vs. AES)
  - ❑ Different trust model (e.g., client vs. server)
- ❑ TEE requires a different set of features
  - ❑ Resource usage (e.g., memory, I/O)
  - ❑ Various constraints (e.g., power, latency)






# Evaluation

- ❑ Security Analysis
  - ❑ Keystone enclave defends various adversary models
- ❑ Modularity Analysis
  - ❑ Keystone supports fine-grained and modular configuration
- ❑ Trusted Computing Base Analysis
  - ❑ Various of real-world applications with less than thousands of LoC
- ❑ Performance Analysis
  - ❑ Security Monitor Overhead
  - ❑ Runtime Overhead
  - ❑ Cost of Memory Protection Mechanisms

# Evaluation

- ☐ Security Analysis
  - ☐ Keystone enclave defends various adversary models
- ☐ Modularity Analysis
  - ☐ Keystone **Please check our paper!** on
- ☐ Trusted Computing Base Analysis
  - ☐ Various of real-world applications with less than thousands of LoC
- ☐ Performance Analysis
  - ☐ Security Monitor Overhead
  - ☐ Runtime Overhead
  - ☐ Cost of Memory Protection Mechanisms

# Runtime Overhead: Memory Management



## □ Torch benchmark

- Unmodified NN inference

## □ Initialization overhead

- Enclave measurement (SHA3)

## □ Execution overhead

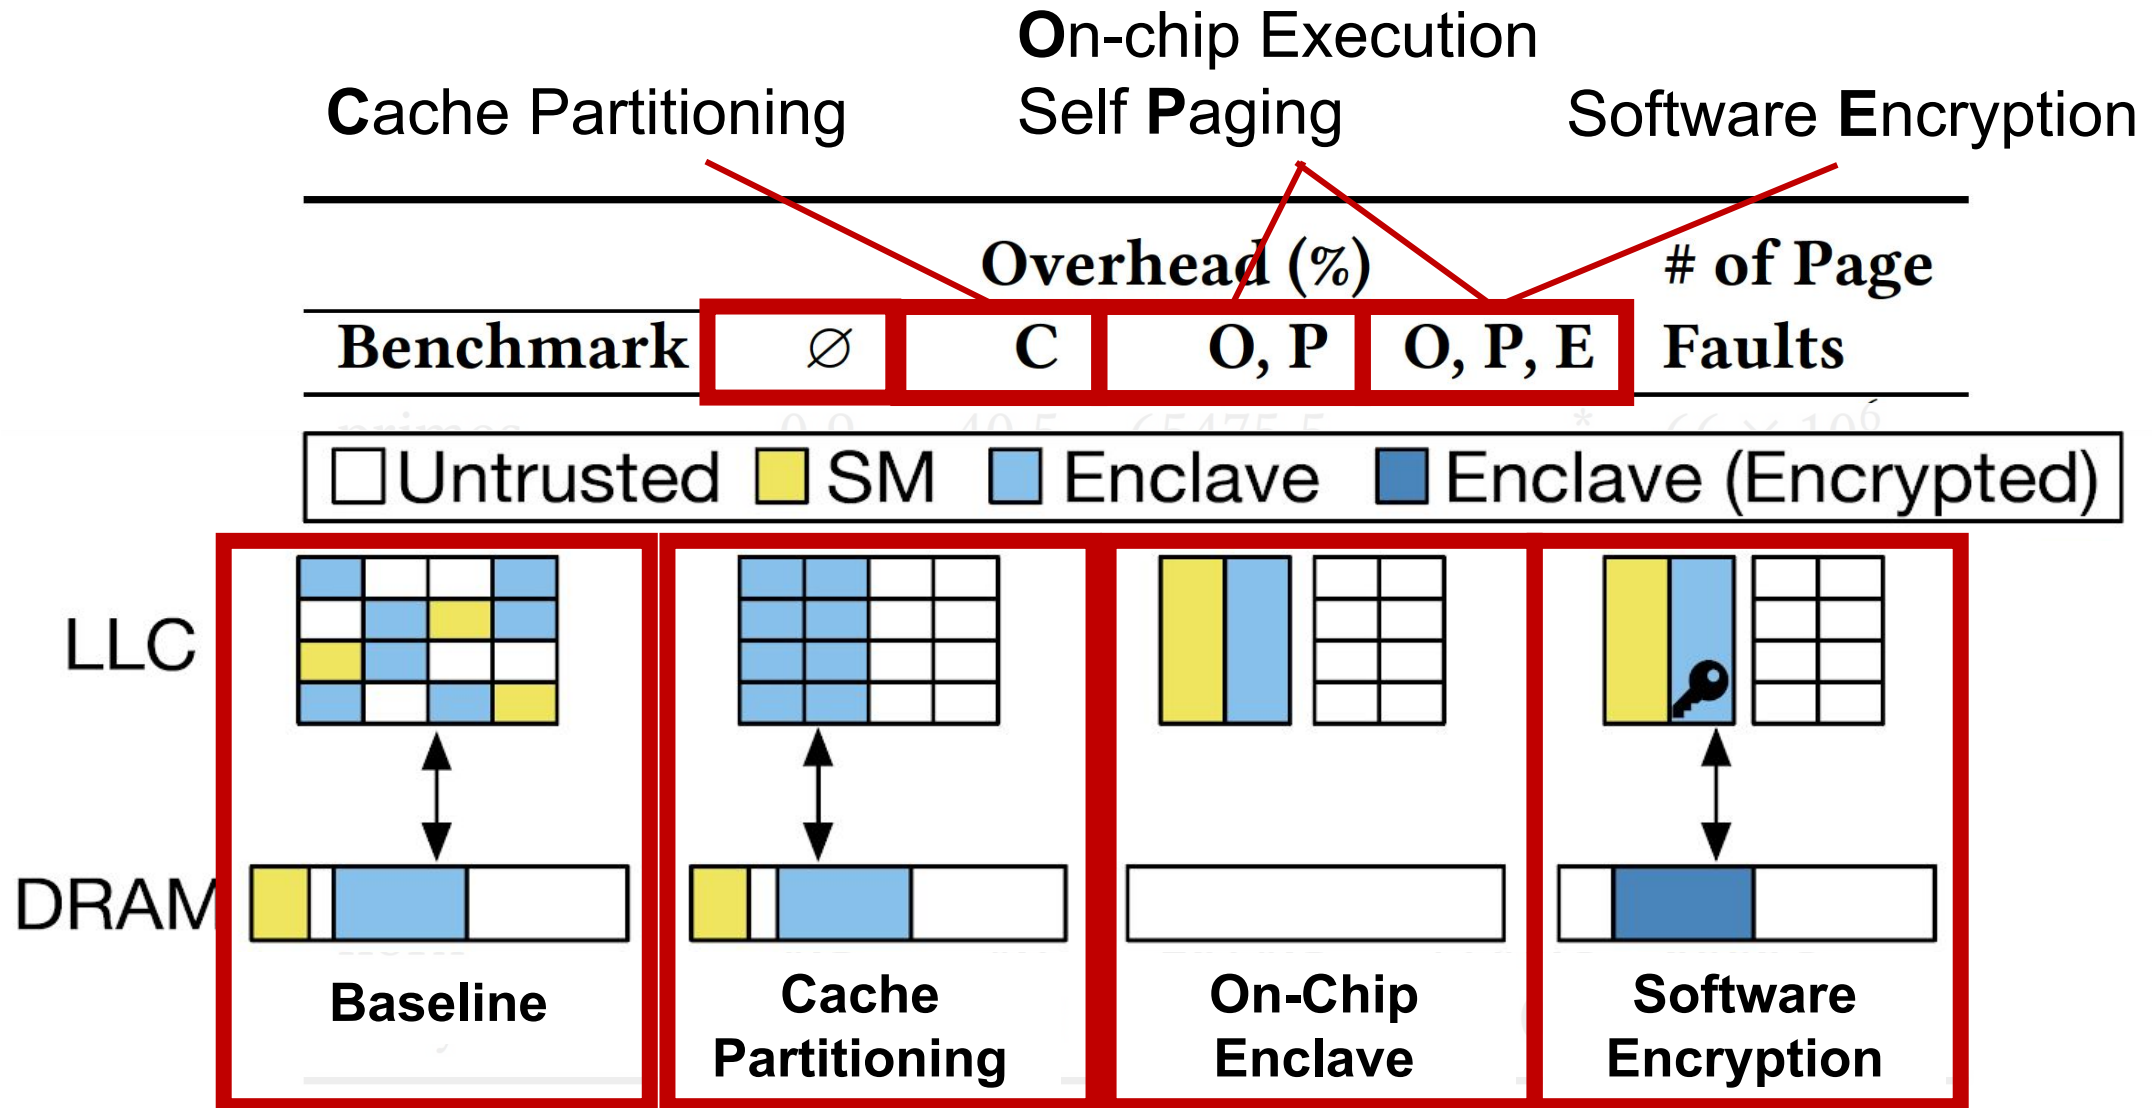
- Min -3.12% (LeNet)

- Max 7.35% (DenseNet)

## □ Dynamic memory resizing

- No noticeable overhead

# Cost of Memory Protection Mechanisms




# Cost of Memory Protection Mechanisms

Cache Partitioning		On-chip Execution Self Paging		Software Encryption	
		Overhead (%)		# of Page Faults	
Benchmark	Ø	C	O, P	O, P, E	Faults
primes	-0.9	40.5	65475.5	*	$66 \times 10^6$
miniz	0.1	128.5	80.2	615.5	18341
aes	-1.1	66.3	1471.0	4552.7	59716
bigint	-0.1	1.6	0.4	12.0	168
qsort	-2.8	-1.3	12446.3	26832.3	285147
sha512	-0.1	0.3	-0.1	-0.2	0
norx	0.1	0.9	2590.1	7966.4	58834
dhrystone	-0.2	0.3	-0.2	0.2	0

# Conclusion

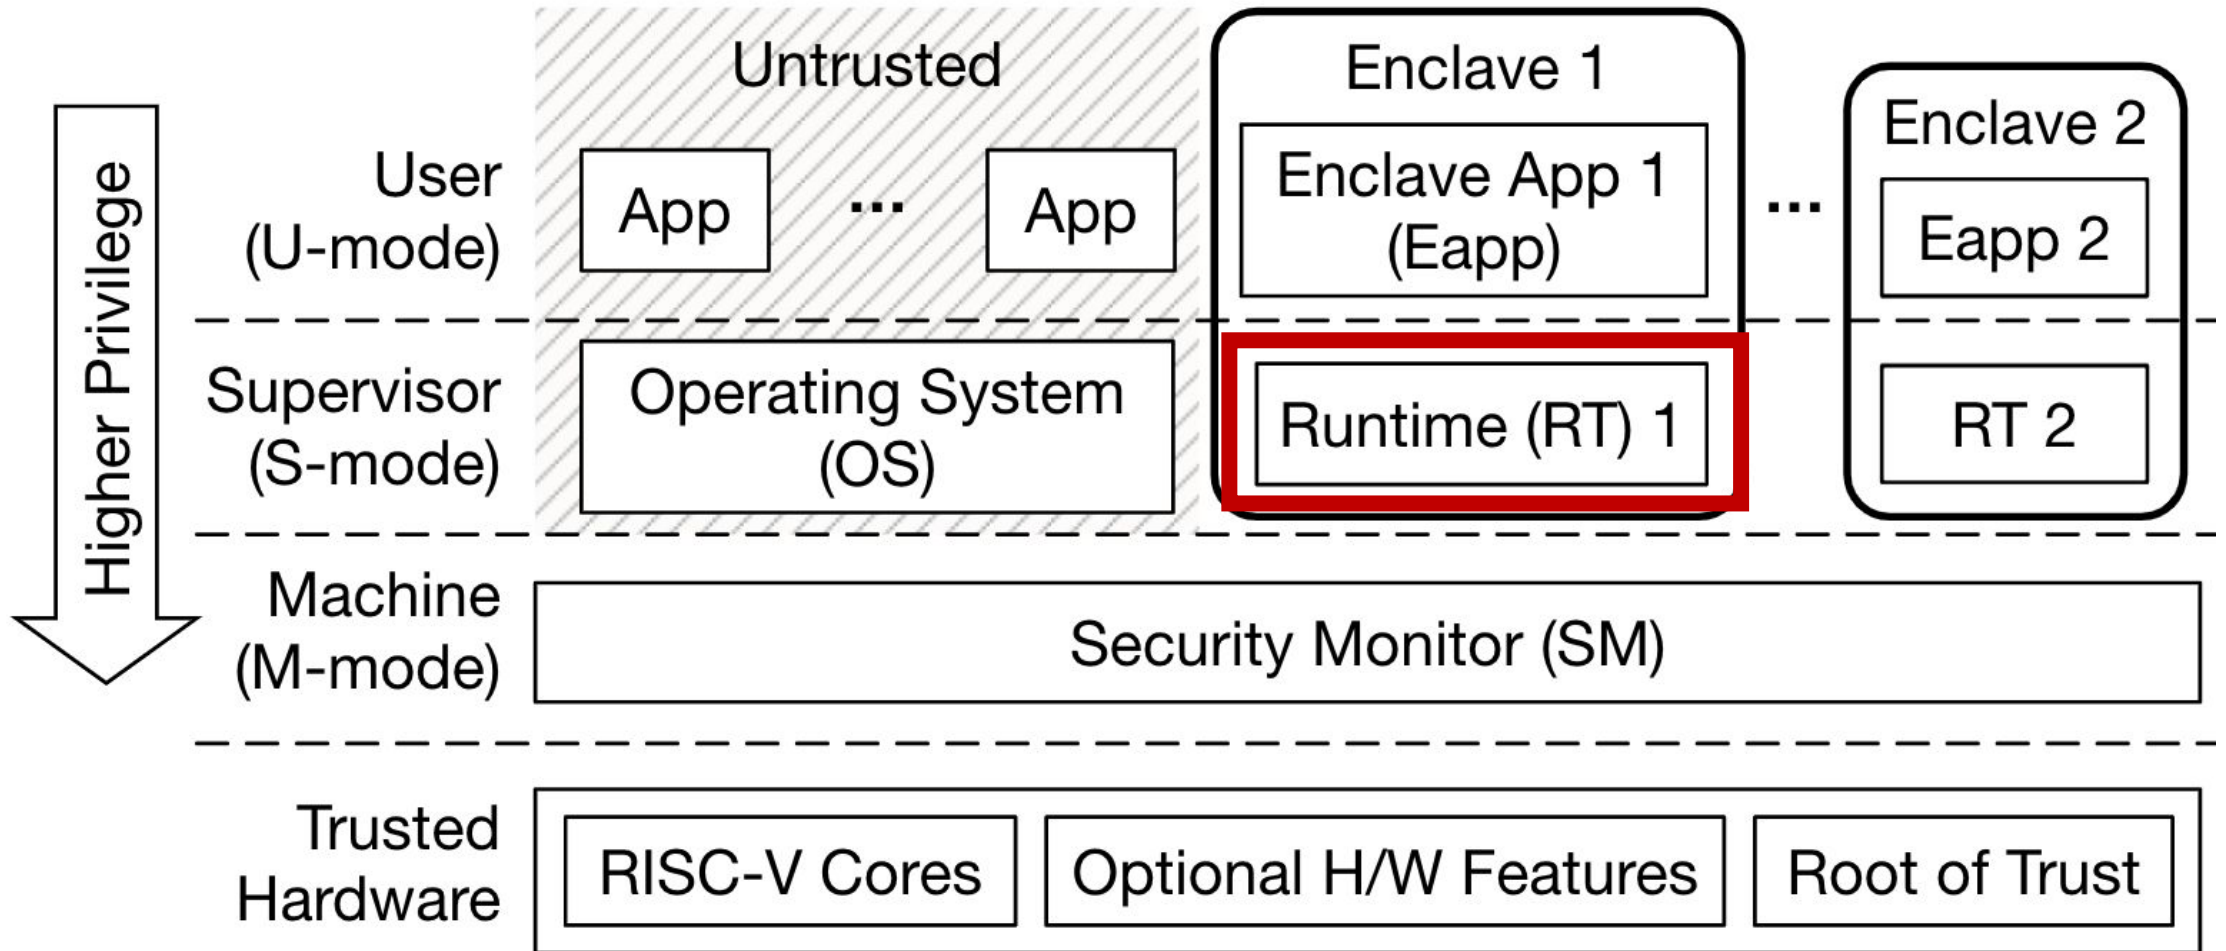
- ❑ Introduced Keystone, a *customizable* framework for TEEs
- ❑ Modular design with fine-grained customizability
- ❑ Useful for building TEEs for different threat models, functionality, and performance requirements
- ❑ Keystone is fully open-source under BSD 3-Clause

# What does Keystone Runtime Do?



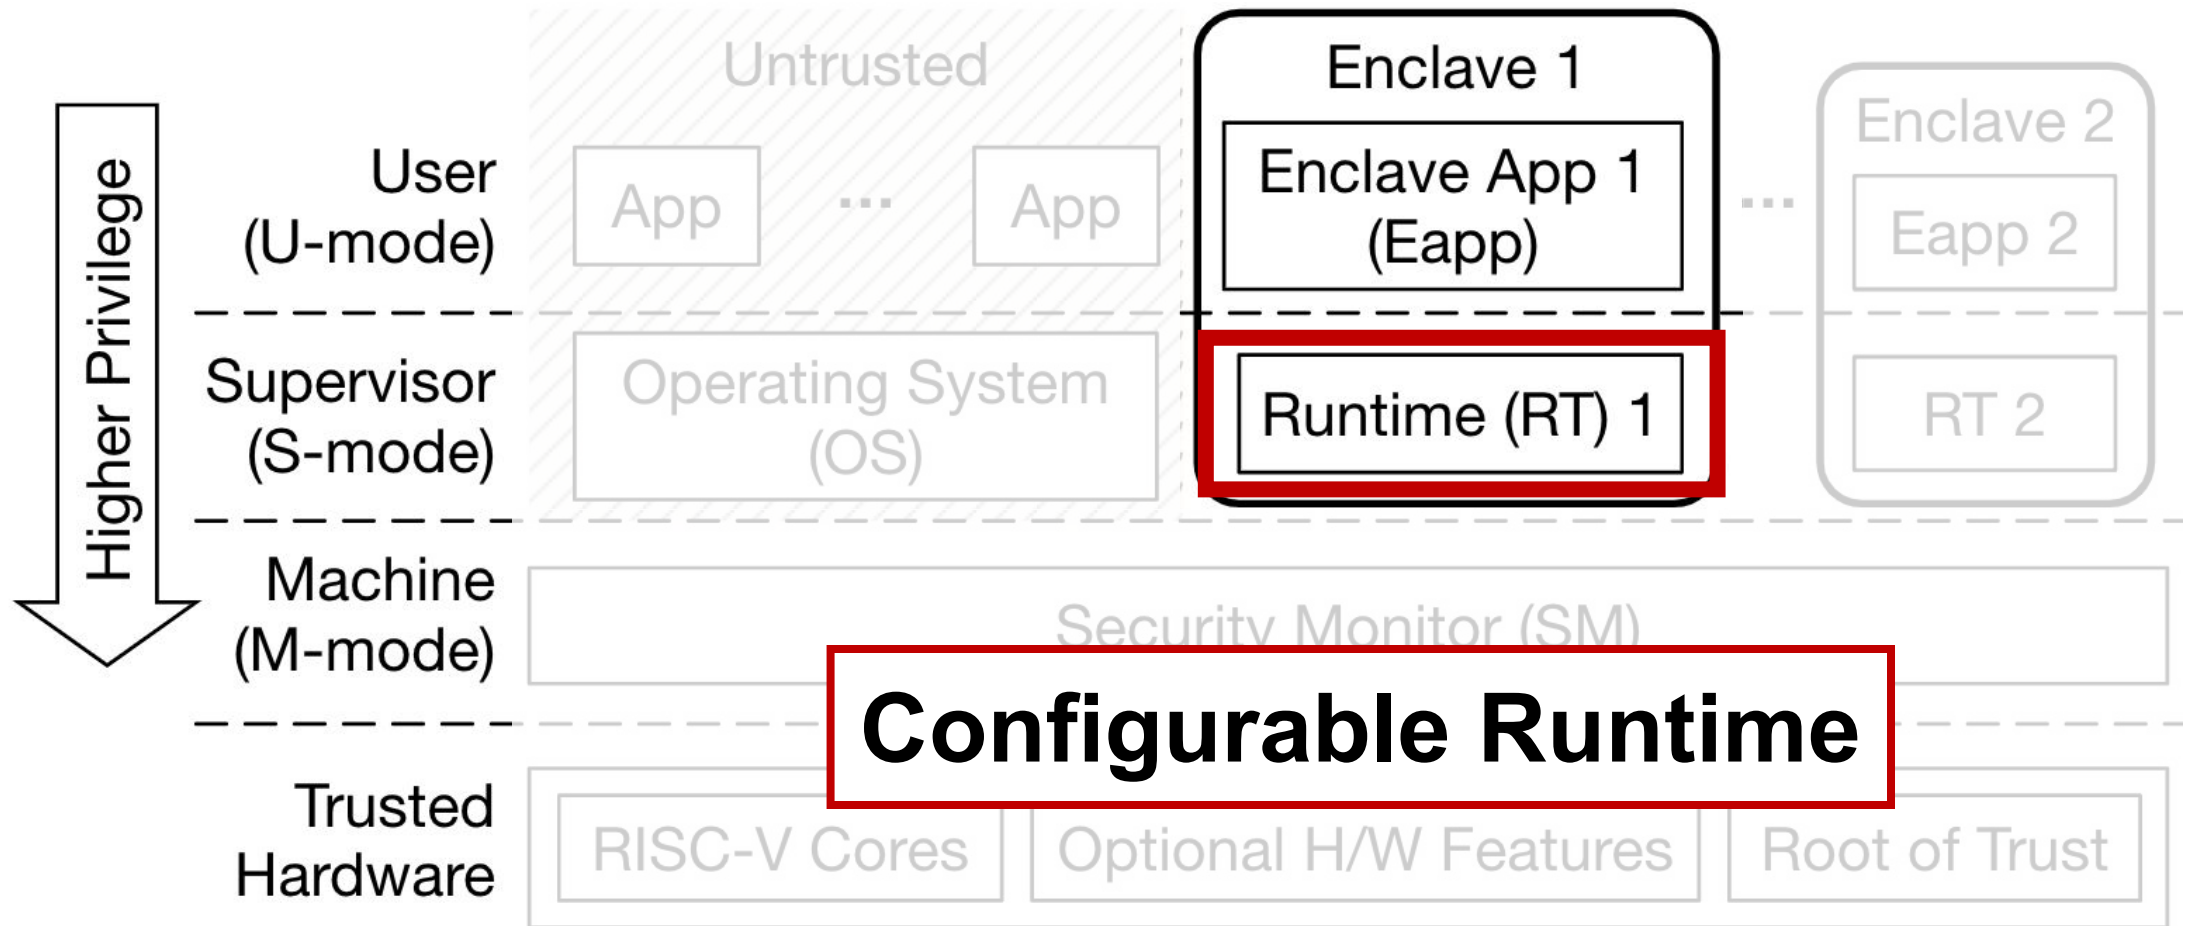
- ☐ A kernel-privileged trusted component for each enclave
  - ☐ Separation of security & functionality
- ☐ Flexible layer of abstraction
  - ☐ Minimal interface & functionality for small TCB (<4,000 LoC)
  - ☐ Fully featured, formally verified kernel (i.e., seL4)
- ☐ Fine-grained customizability for enclaves
  - ☐ Memory management: free memory, self-paging, memory encryption
  - ☐ Functionality: libraries (e.g., libc, musl-libc) and system calls

# What does Keystone Runtime Do?

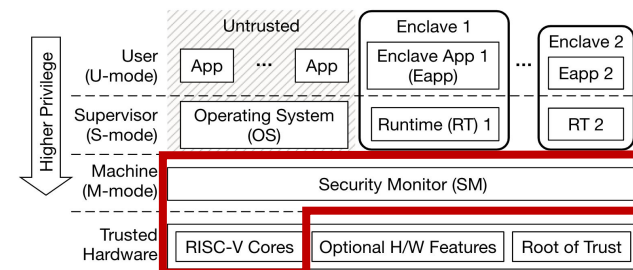




# Keystone Architecture and Trust Model





# Memory Isolation via RISC-V PMP



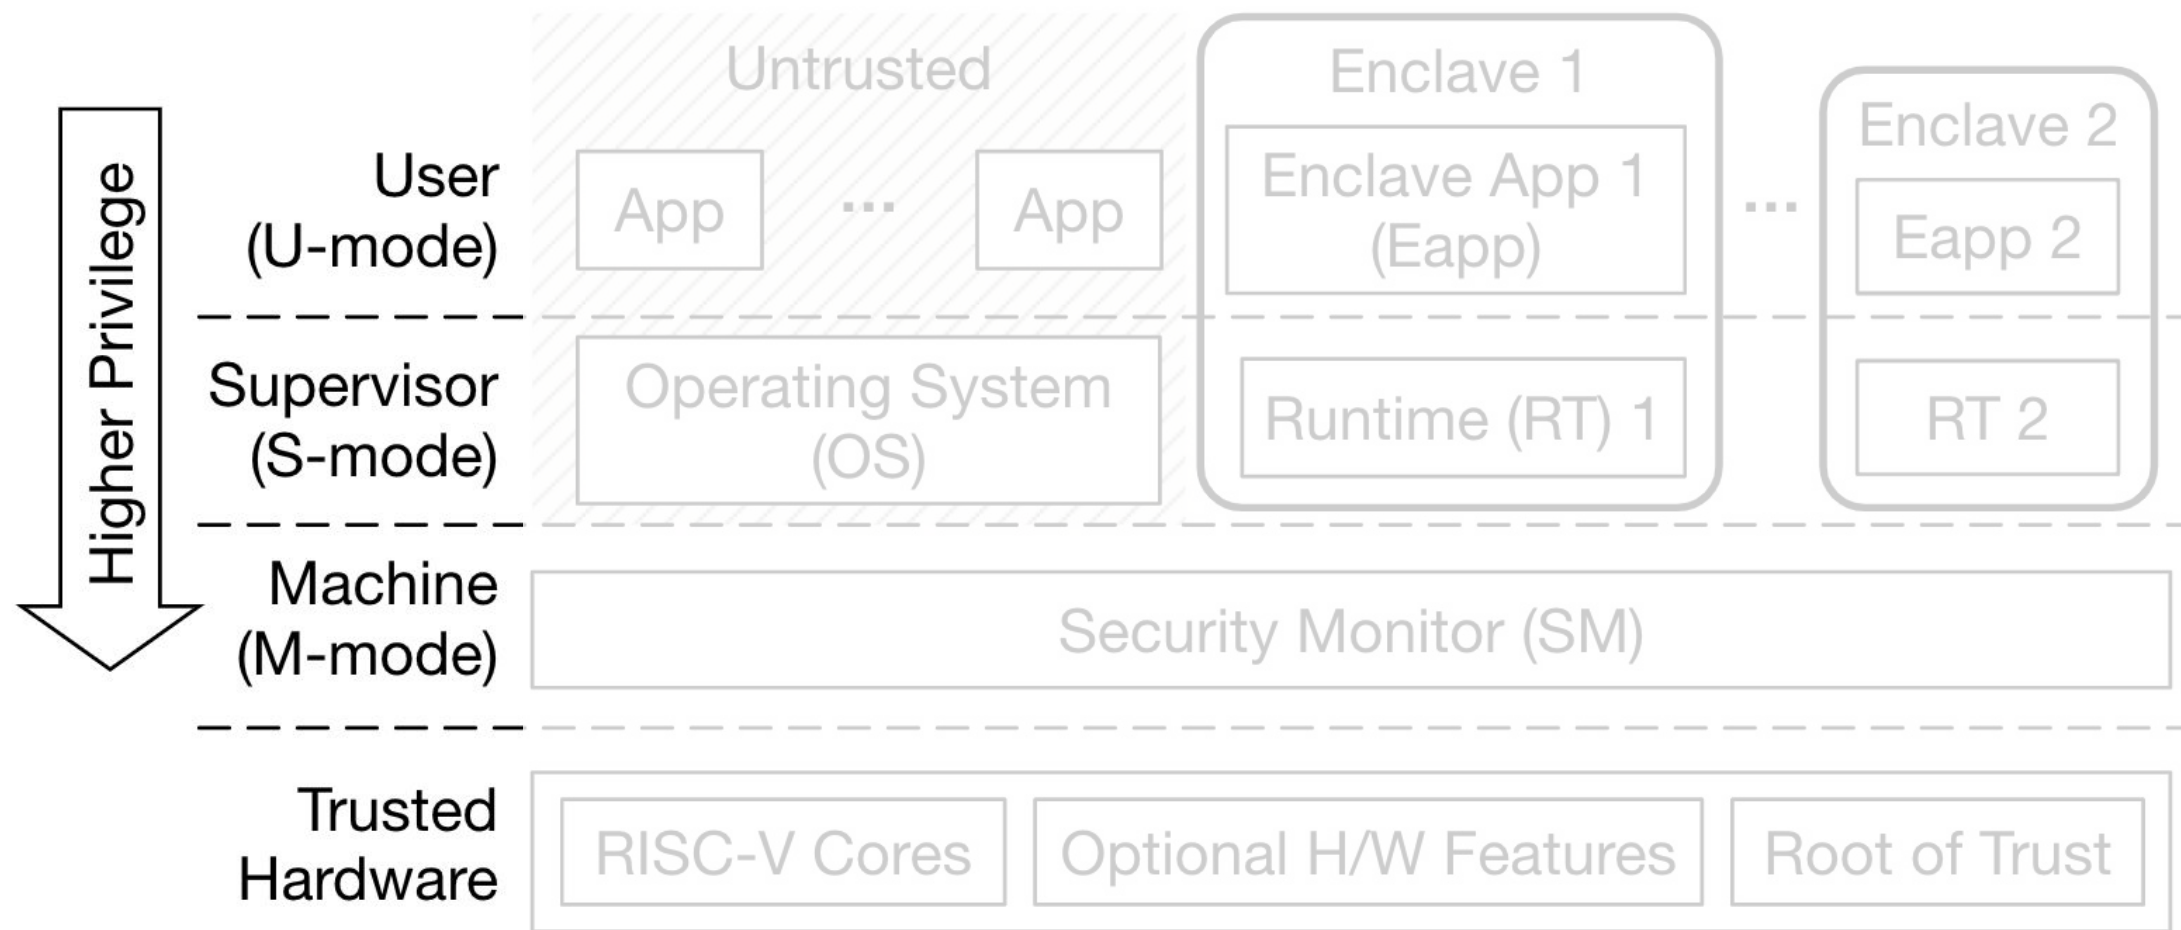
## PMP registers



## pmpcfg perm.

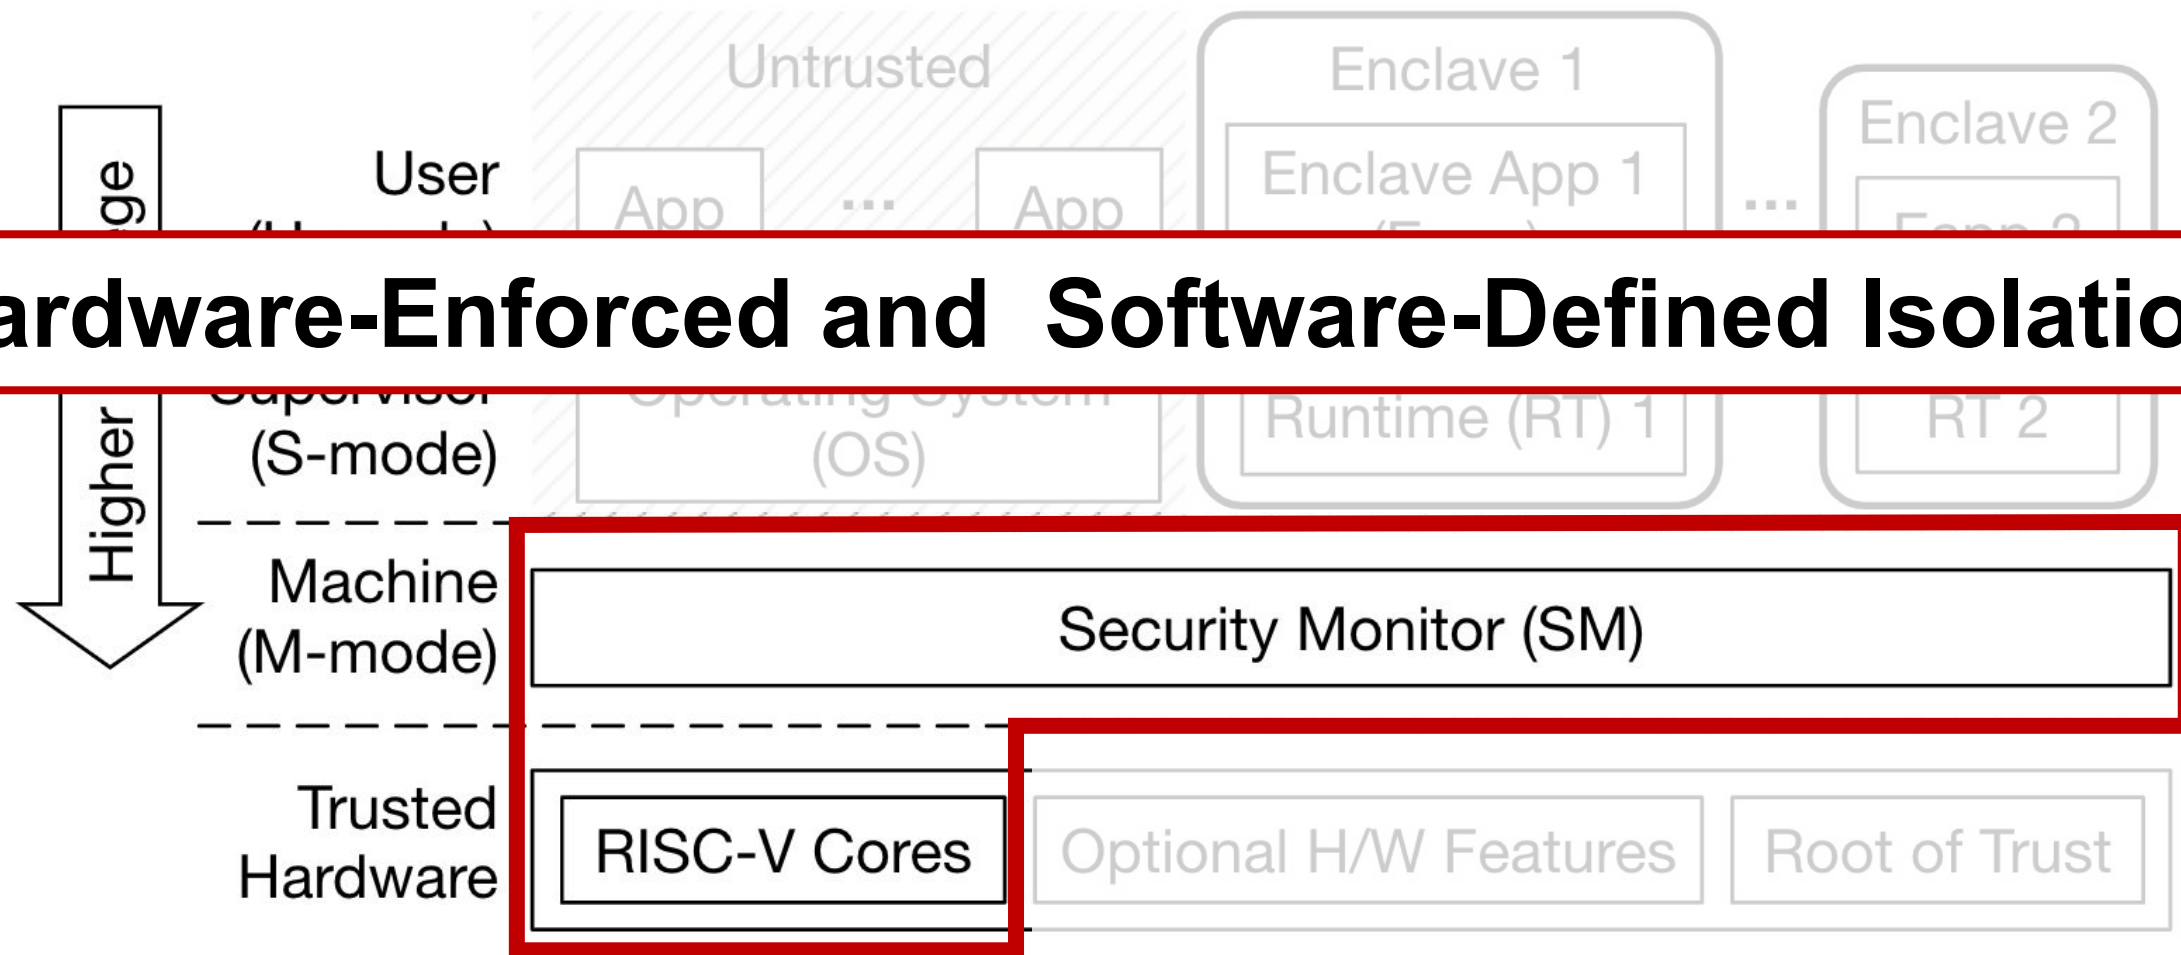


# Keystone Architecture and Trust Model

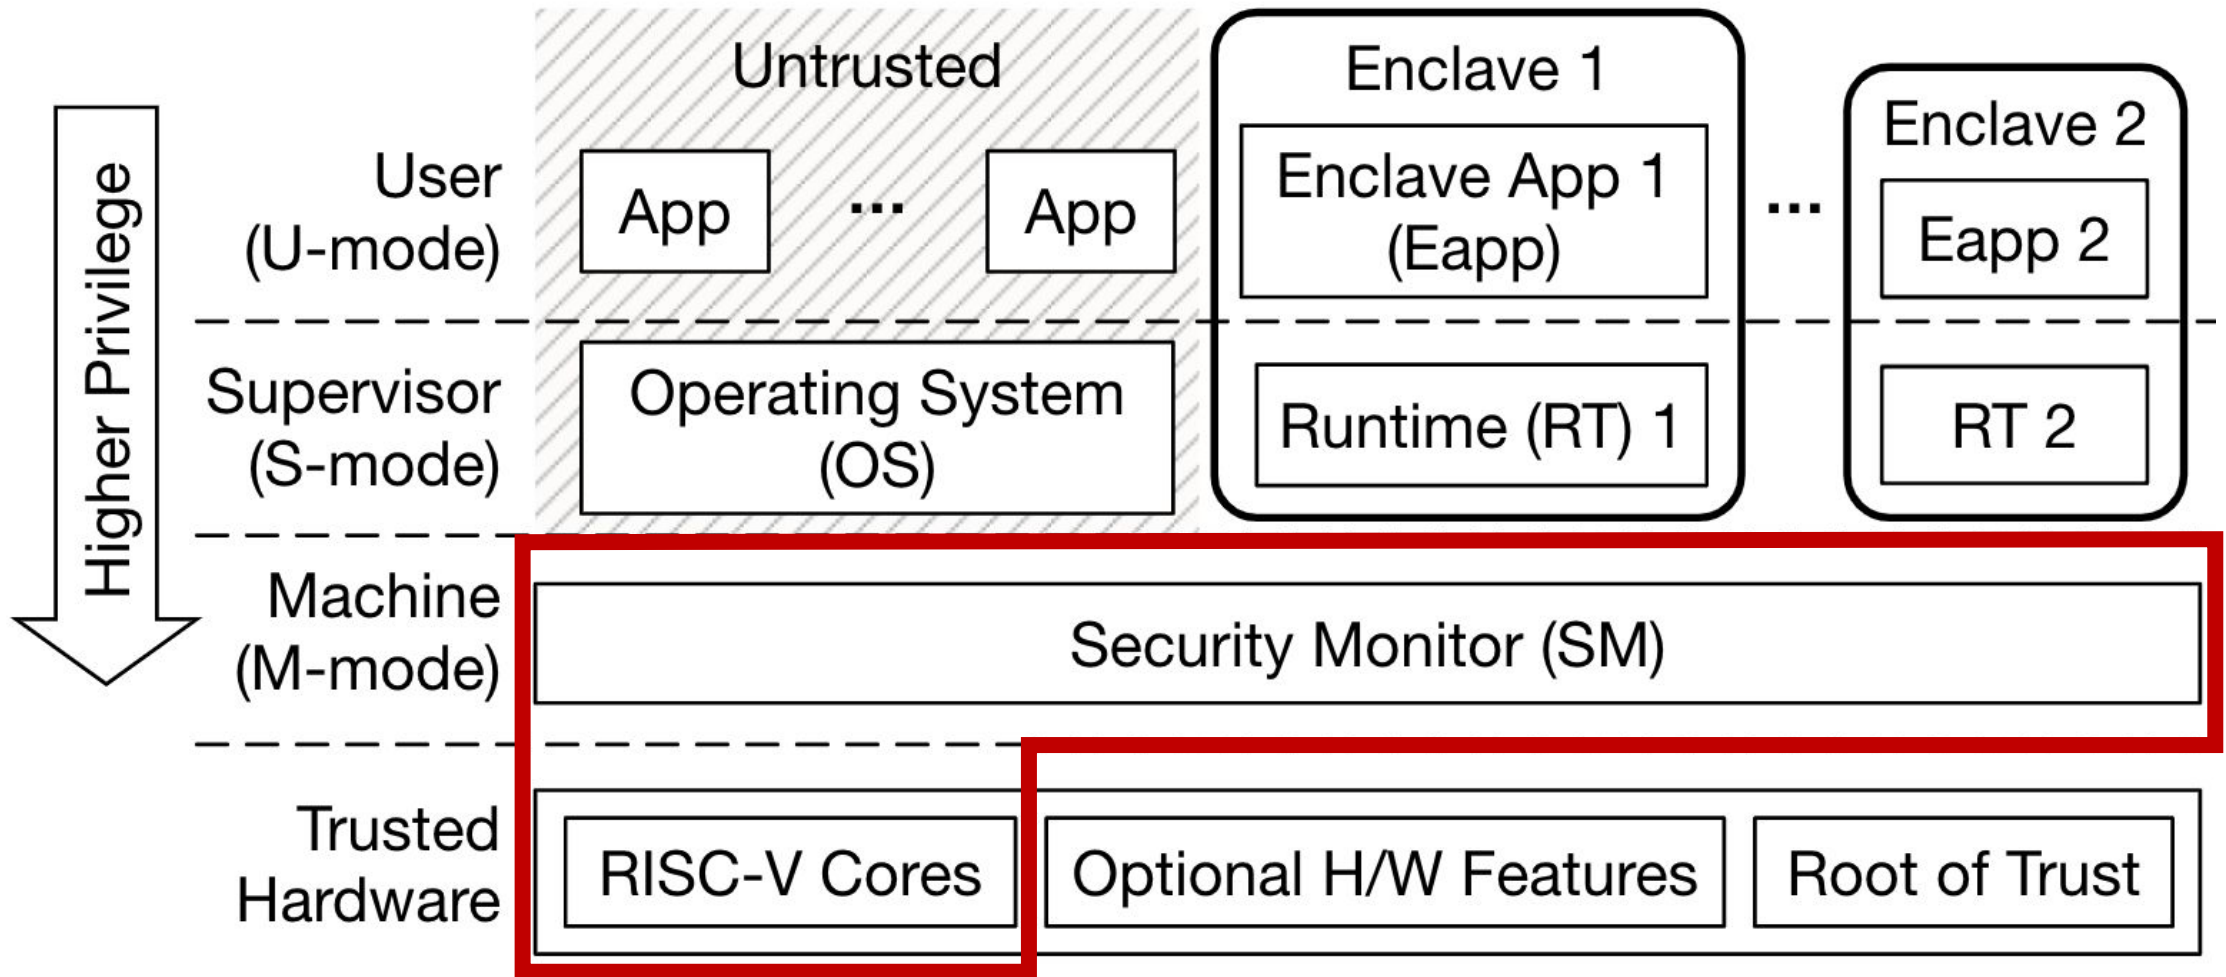


# Keystone Architecture and Trust Model

## Hardware-Enforced and Software-Defined Isolation

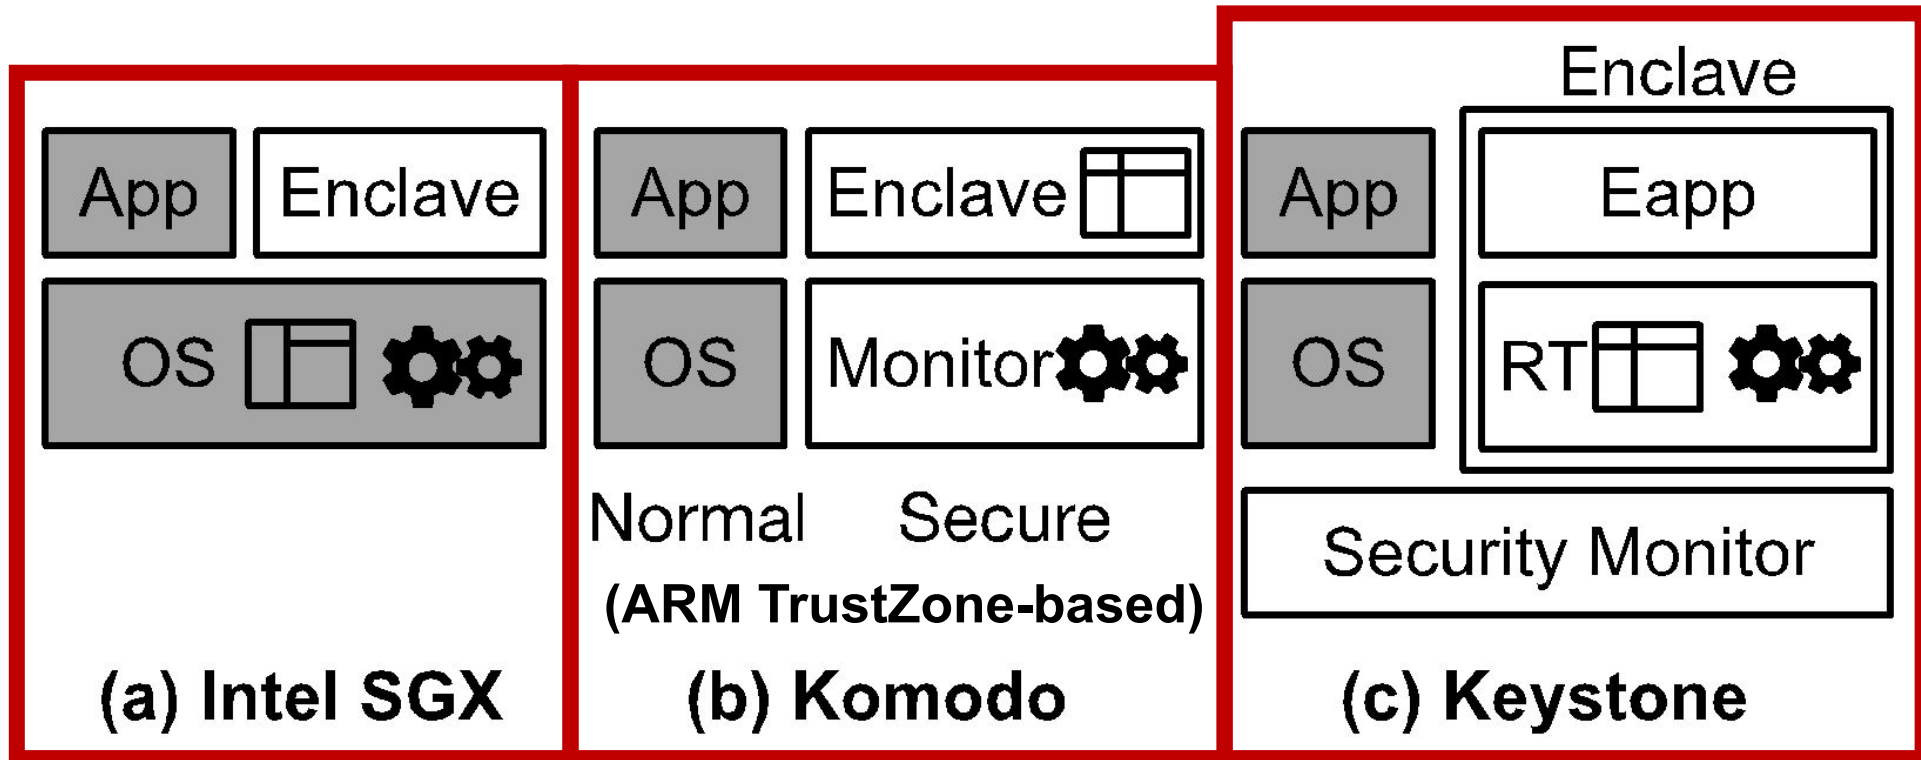


# Memory Isolation via RISC-V PMP



# Memory Management in Keystone

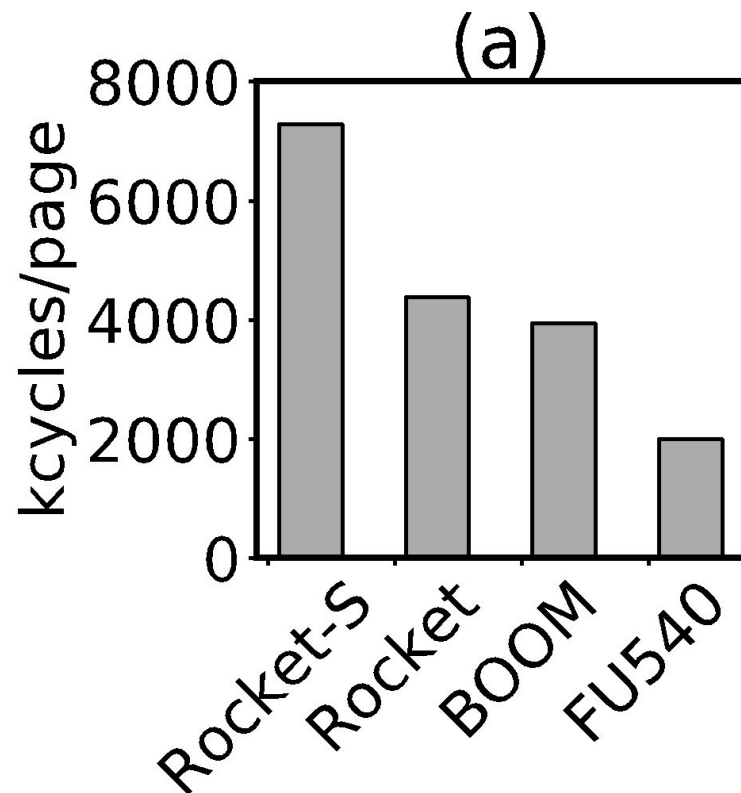
- (1) Does the host share virtual addresses with the enclave?
- (2) Who owns the memory management unit (MMU)?



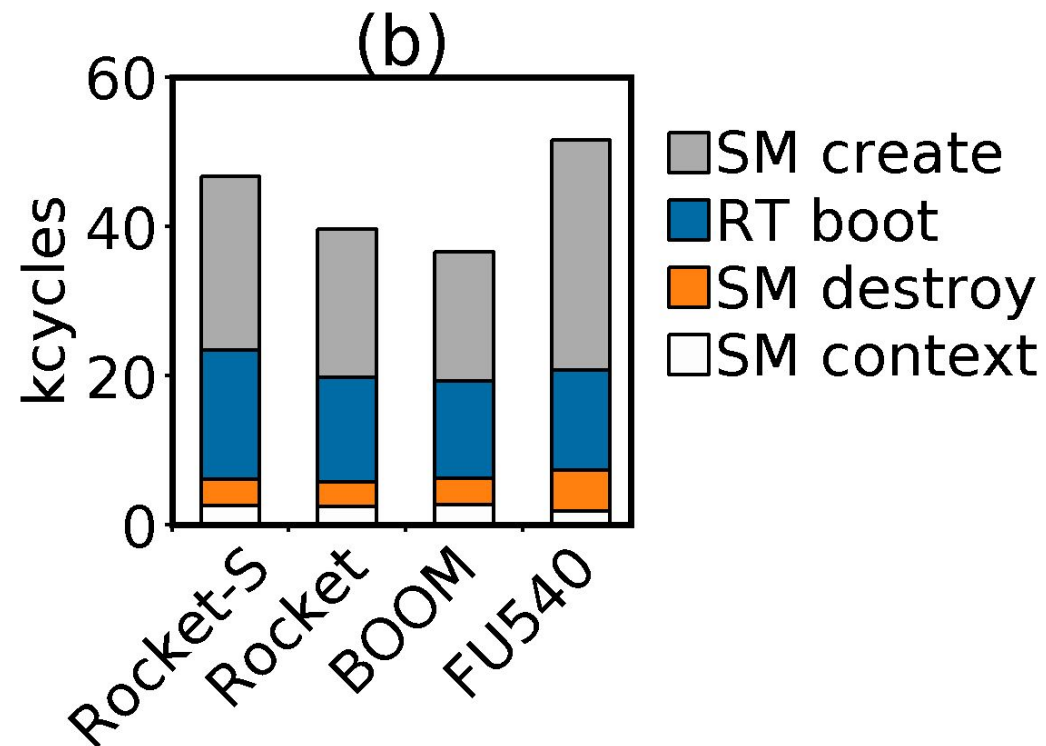
- ☐ Flexible resource management (e.g., dynamic memory resizing)
- ☐ Flexible memory protection mechanisms

# Security Monitor Overhead

## Enclave Measurement



## Security Monitor Overhead



- ❑ Long initialization due to measurement (SHA3)
- ❑ No execution overhead in CPU benchmarks (CoreMark, beebs)



# Contributions


- ❑ Provide a **common base** for diverse TEEs
  - ❑ Security monitor <sup>[1,2]</sup>: programmable trusted layer *below kernel/hypervisor*
  - ❑ *Hardware-enforced* memory/context isolation and attestation
- ❑ A software framework for **customizable TEE**
  - ❑ Separation of security and functionality (e.g., resource mgmt.)
  - ❑ Fine-grained configuration of modular extensions
- ❑ Benchmarking & real-world applications
  - ❑ Overhead of various operations (e.g., enclave creation, I/O)
  - ❑ Performance trade-offs for various defenses (e.g., cache partitioning)
- ❑ An open-source, full-stack implementation for further research

[1] Costan et al., USENIX Sec'16


[2] Ferraiuolo et al., SOSP'17




# Isolation with RISC-V PMP




# Isolation with RISC-V PMP



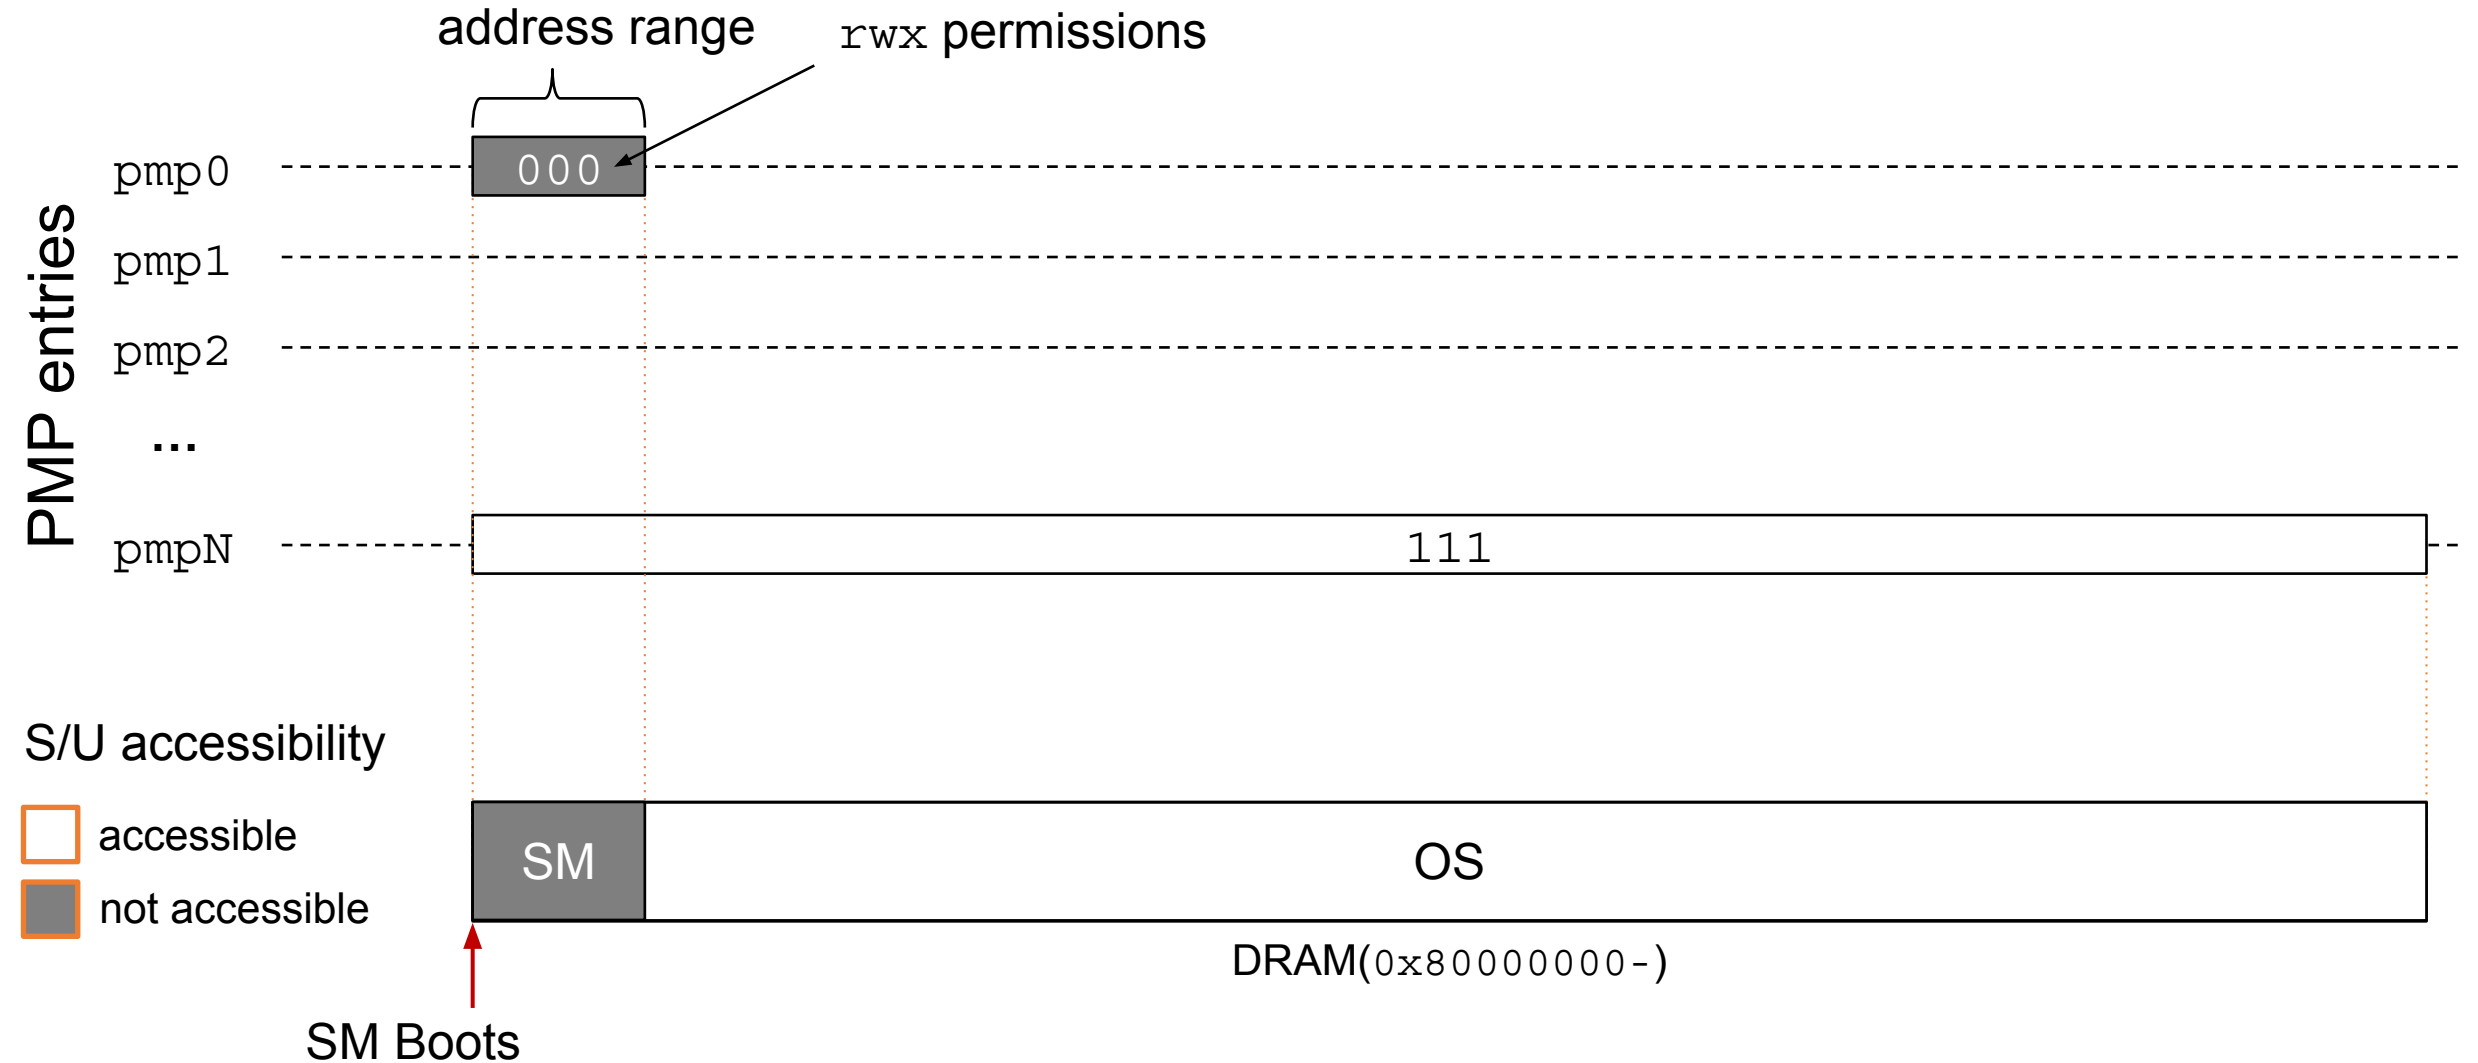
# Isolation with RISC-V PMP



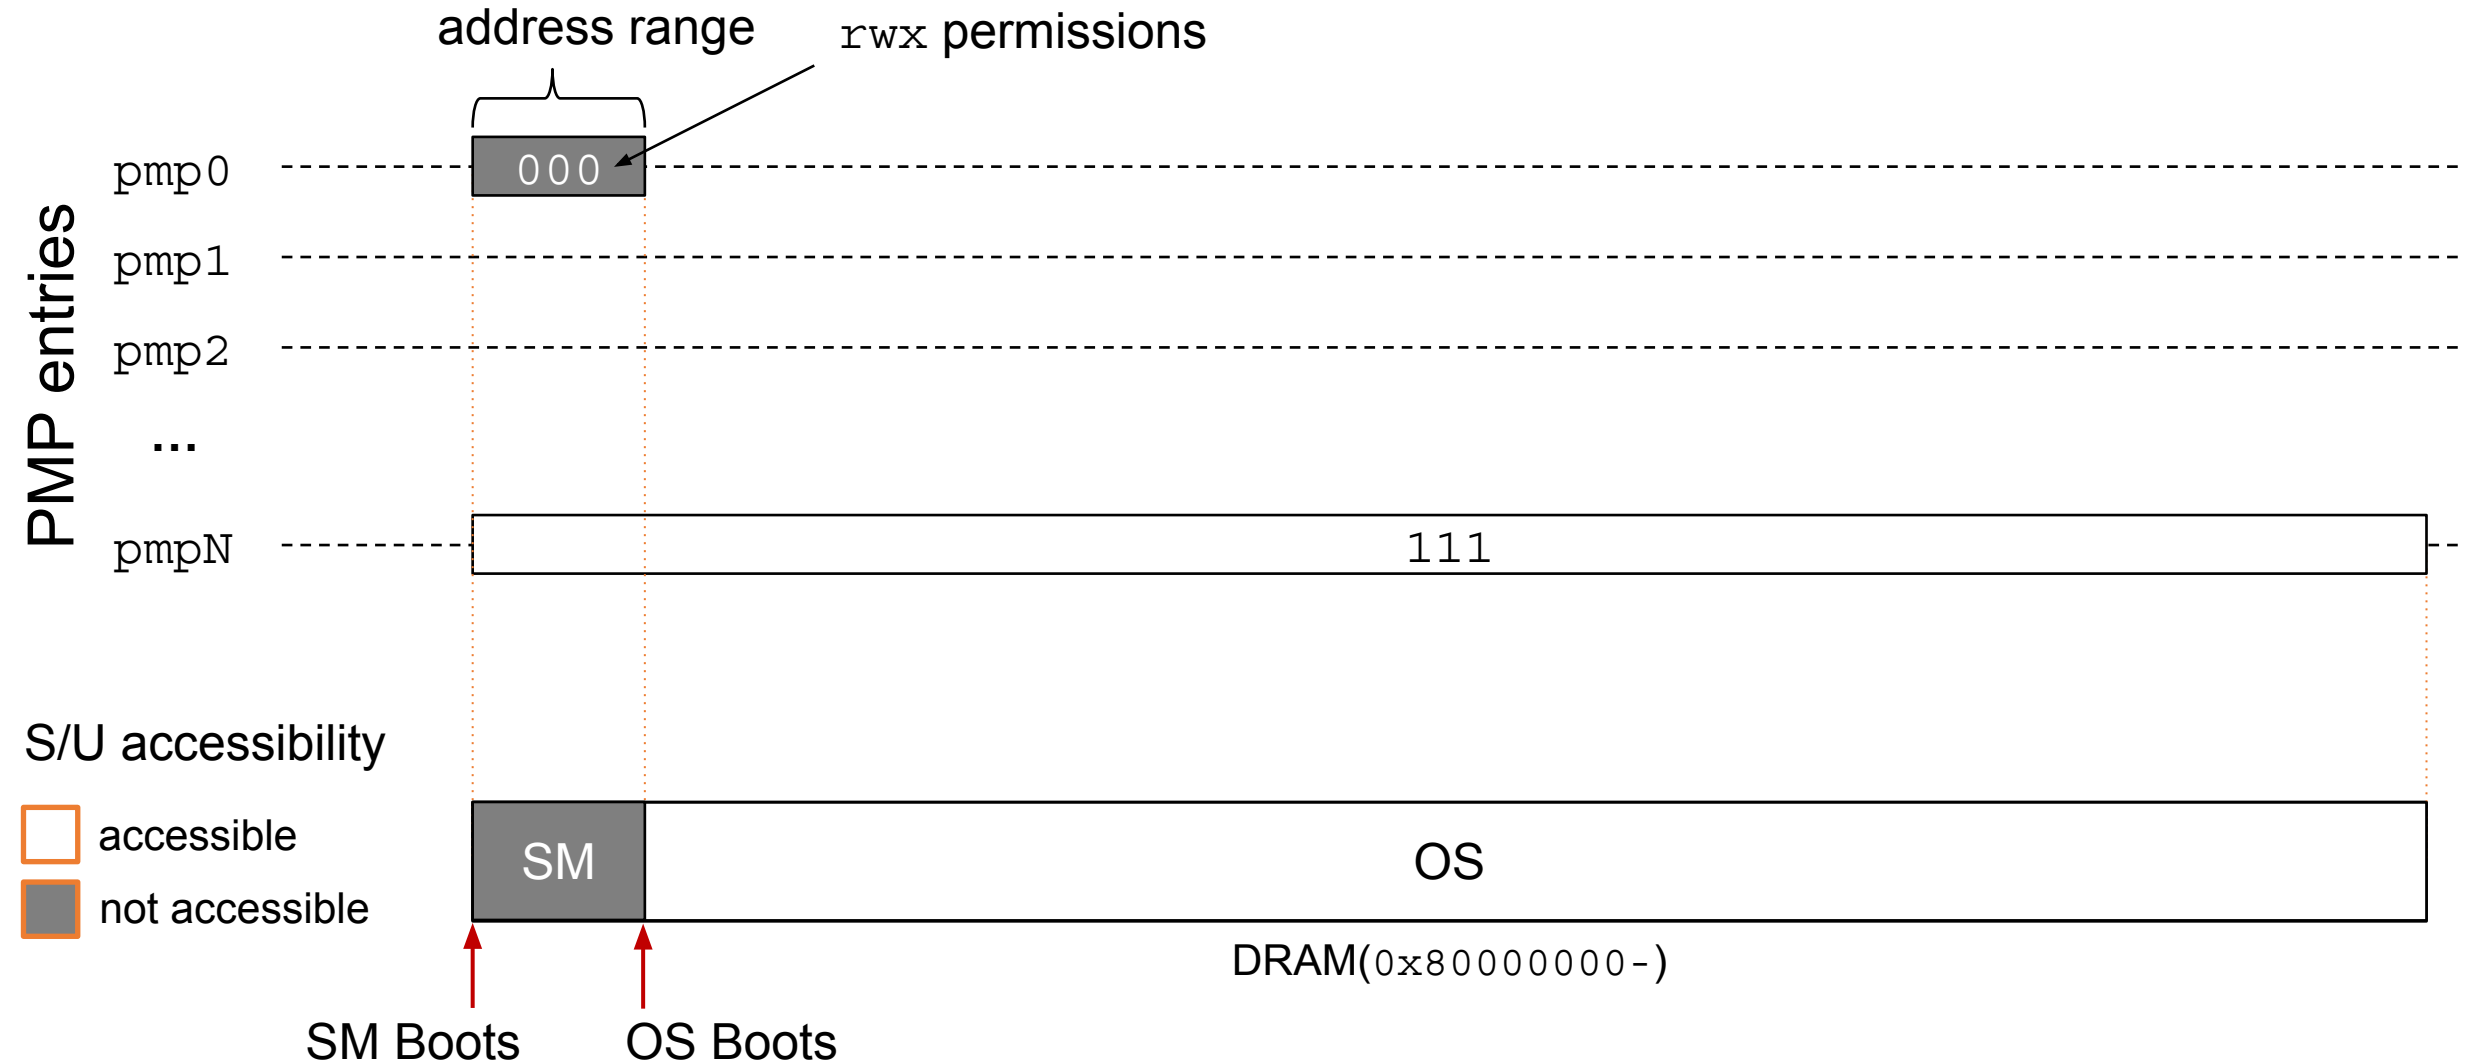
# Isolation with RISC-V PMP



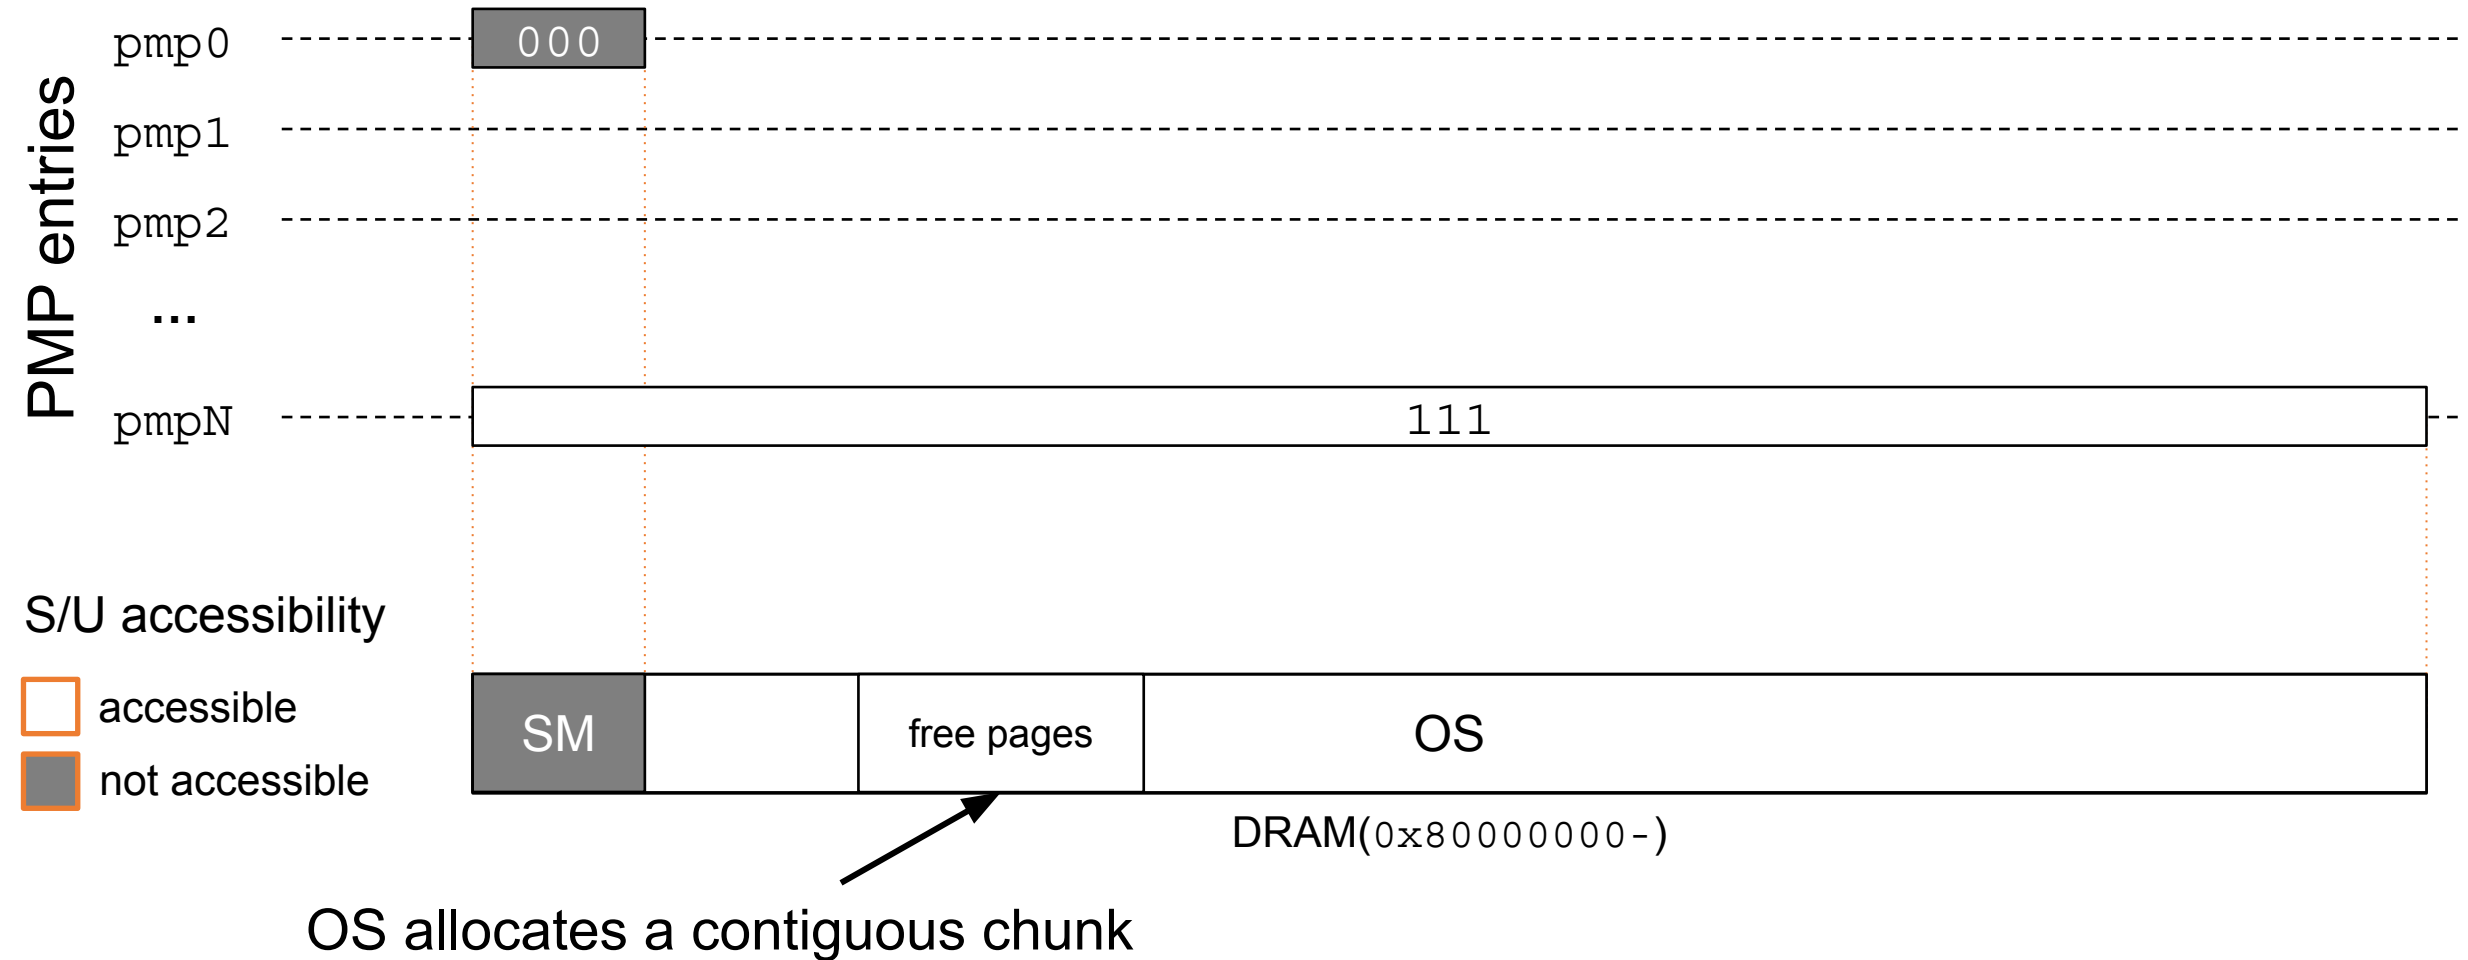
# Isolation with RISC-V PMP




# Isolation with RISC-V PMP



# Creating an Isolated Enclave




# Creating an Isolated Enclave






# Creating an Isolated Enclave




# Creating an Isolated Enclave

OS can ask the SM to create multiple enclaves




# Executing an Enclave



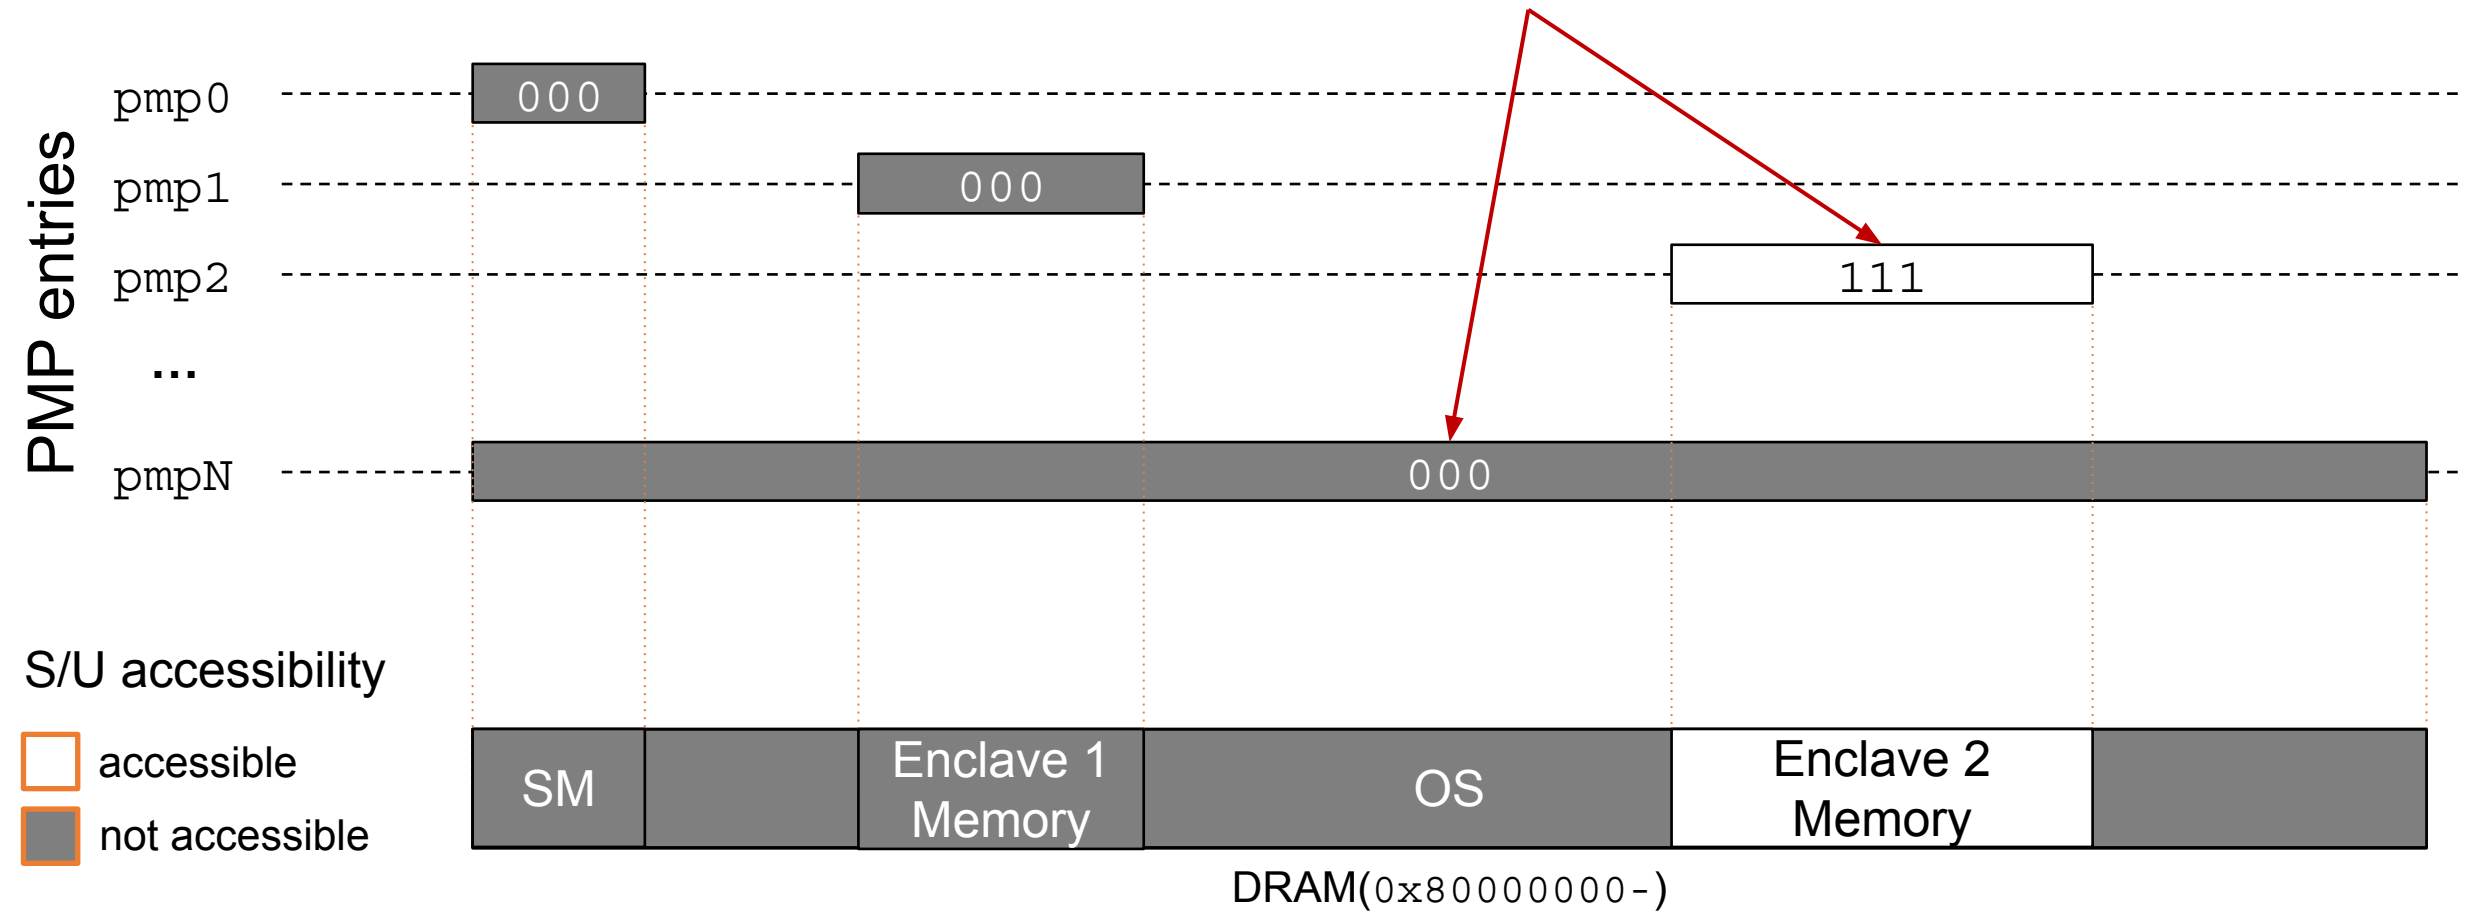
# Executing an Enclave

For Enclave 2 SM sets rwx for pmp2 and sets - - - for pmpN




# Executing an Enclave

For Enclave 2 SM sets rwx for pmp2 and sets - - - for pmpN




# Exit an Enclave

Switch back to defaults for the OS

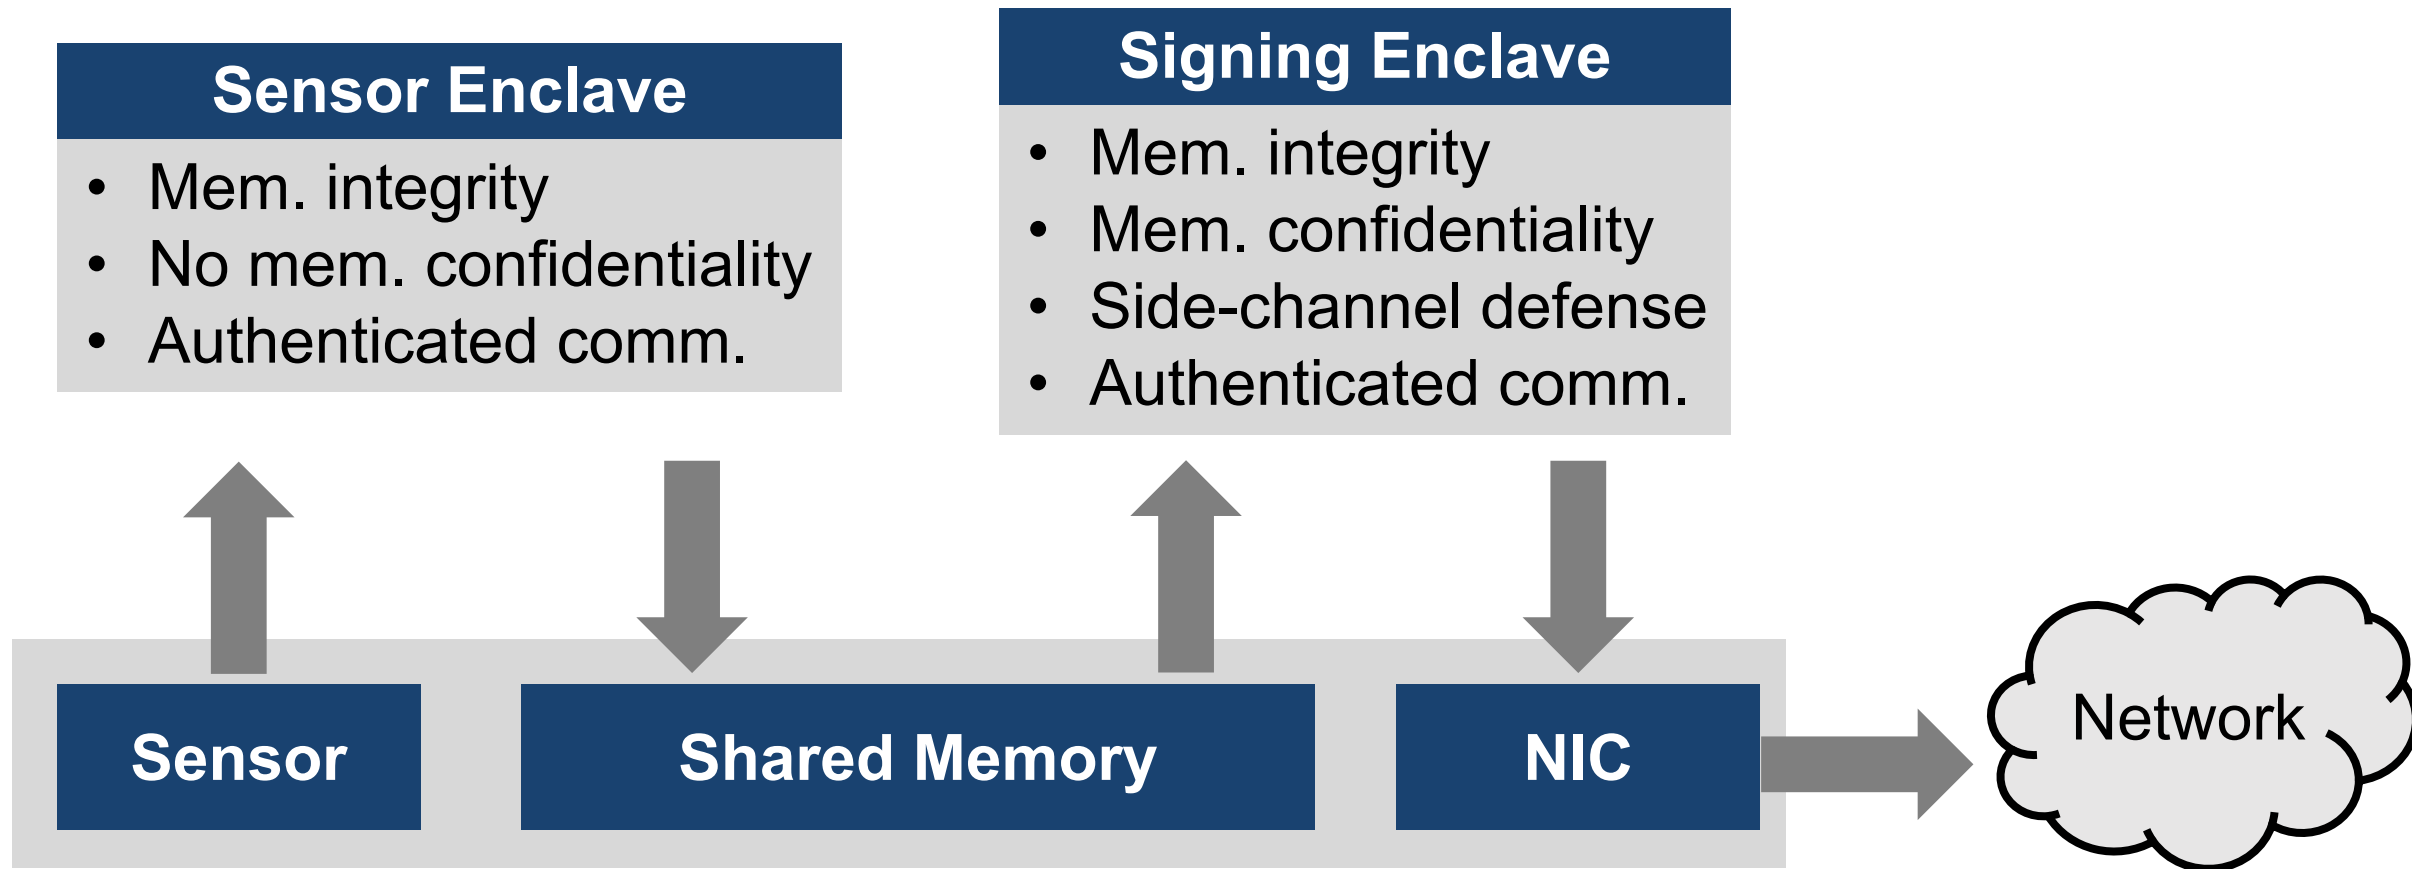


# Destroying an Enclave



# Varying Threat Models in the Same Platform

e.g., IoT Sensor Platform

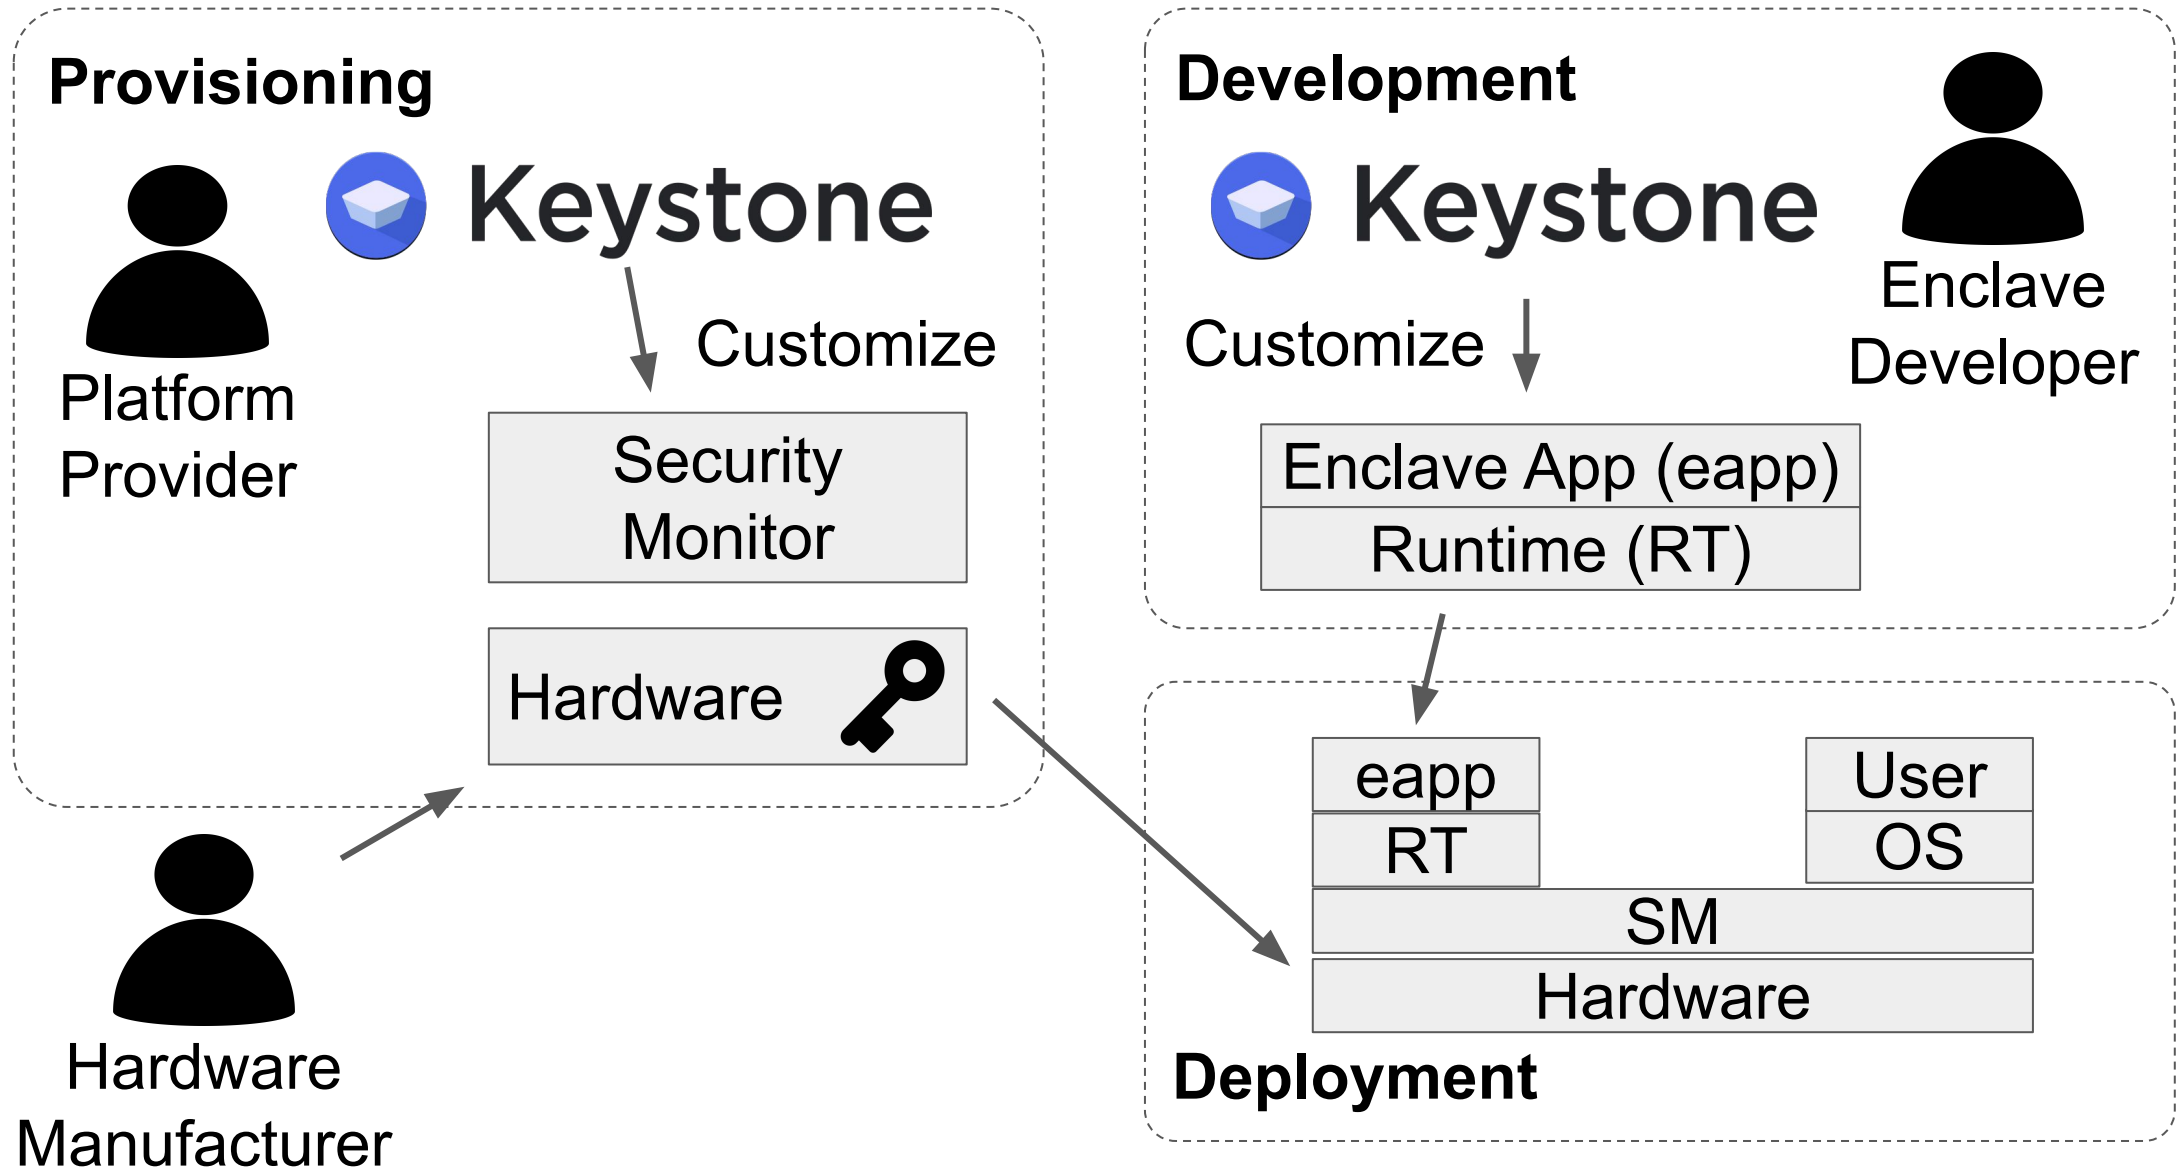





# Why TEEs are Promising for Security?

- ❑ Increasing needs for secure computation
  - ❑ Big data and machine learning (e.g., cooperative learning)
  - ❑ Trends in IoT, mobile, and cloud computing
  - ❑ Requirements in new applications (e.g., blockchain, contact tracing)
- ❑ Security in system software is getting harder
  - ❑ Increasing attack surface in the system software (e.g., Linux kernel)
  - ❑ SW boundaries are often broken by SW/HW attackers
- ❑ TEE as a cornerstone for secure computation
  - ❑ Minimizing trusted computing base (TCB)
  - ❑ Efficient HW-enforced isolation and authentication

# How TEE Customization Work?



# TEE Threat Model Evolves Over Time



Unknown threat is newly discovered (e.g., vulnerability in speculative execution)

Known threats become substantial (e.g., cache side-channel attacks)

# Inflexible Design and Implementation

- ❑ TEEs in commercial hardware: Intel SGX, ARM TZ, AMD SEV
- ❑ Designs and threat models depend too much on their business
  - ❑ Intel SGX – small server/desktop apps (e.g., DRM, cryptography, etc)
  - ❑ ARM TZ – vendor-provisioned mobile apps (e.g., fingerprint, ledger)
  - ❑ AMD SEV – full VM isolation only (targeting cloud market?)
- ❑ Implemented on closed-source hardware
  - ❑ Slow iteration dictated by a company; researchers can't step forward
  - ❑ Any additional features/defenses need significant workaround

# Vendor-Locked Threat Models

ISA/Arch	System	SW Attacks	HW Attacks	Side Channel	Controlled Channel
Intel	SGX	●	●	○	○
	Haven	●	●	○	○
	Graphene	●	●	○	○
ARM	TrustZone	●	○	○	○
	Komodo	●	●	○	●
	OPTEE	●	●	○	○
AMD	SEV	●	●	○	○
	SEV-ES	●	●	○	○
RISC-V	Sanctum	●	○	●	●
	Keystone	●	●	●	●

# Runtime Overhead: Untrusted I/O

